

CNN BASED STAY AWAKE STAY ALIVE SYSTEM FOR ACCIDENT PREVENTION

A Project Report submitted

in partial fulfillment for the award of the Degree of

**Bachelor of Technology
in
Computer Science and Engineering
by**

B. Deepak (U20CS166)

B. Naga Sudhakar (U20CS156)

Ch. Rakesh (U20CS174)

D. Tirumalesh (U20CS214)

Under the guidance of

Dr P. Vasuki , M.E., Ph.D




**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
SCHOOL OF COMPUTING**

BHARATH INSTITUTE OF HIGHER EDUCATION AND RESEARCH

(Deemed to be University Estd u/s 3 of UGC Act, 1956)

CHENNAI 600073, TAMILNADU, INDIA

April, 2024



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

BONAFIDE CERTIFICATE

This is to Certify that this Project Report Titled “CNN BASED STAY AWAKE STAY ALIVE SYSTEM FOR ACCIDENT PREVENTION” is the Bonafide Work of B. Deepak (U20CS166), B. Naga Sudhakar (U20CS156), CH. Rakesh (U20CS174) and D. Tirumalesh (U20CS214) of Final Year B.Tech. (CSE) who carried out the project work under my supervision Certified further, that to the best of my knowledge the work reported here in does not form part of any other project report or dissertation on basis of which a degree or award conferred on an earlier occasion by any other candidate.

PROJECT GUIDE

Dr. P. Vasuki M.E., Ph.D

Professor

Department of CSE

BIHER

HEAD OF THE DEPARTMENT

Dr. S. Maruthuperumal

Professor

Department of CSE

BIHER

Submitted for the Project Viva-Voice held on

INTERNAL EXAMINER

EXTERNAL EXAMINER

DECLARATION

We declare that this project report titled **CNN BASED STAY AWAKE STAY ALIVE SYSTEM FOR ACCIDENT PREVENTION** submitted in partial fulfillment of the degree of **B. Tech in Computer Science and Engineering** is a record of original work carried out by us under the supervision of **Dr P. Vasuki**, and has not formed the basis for the award of any other degree or diploma, in this or any other Institution or University. In keeping with the ethical practice in reporting scientific information, due acknowledgments have been made wherever the findings of others have been cited.

B. Deepak
(U20CS166)

B. Naga Sudhakar
(U20CS156)

Ch. Rakesh
(U20CS174)

D. Tirumalesh
(U20CS214)

Chennai

Date:

ACKNOWLEDGEMENTS

We express our sincere thanks to our beloved Honorable Chairman **Dr.S.Jagathrakshakan, M.P.**, for continuous and constant encouragement in all academic activities.

We express our deepest gratitude to our beloved President **Dr. J. Sundeep Aanand**, and Managing Director **Dr. E. Swetha Sundeep Aanand** for providing us the necessary facilities to complete our project.

We take great pleasure in expressing sincere thanks to **Dr. K. Vijaya Baskar Raju** Pro Chancellor, **Dr. M. Sundararajan** Vice Chancellor (i/c), **Dr. S. Bhuminathan** Registrar and **Dr. R. Hariprakash** Additional Registrar, **Dr.M. Sundararaj** Dean Academics for moldings our thoughts to complete our project.

We thank our **Dr. S. Neduncheliyan** Dean, School of Computing for his encouragement and the valuable guidance.

We record indebtedness to our Head, **Dr. S. Maruthuperumal**, Department of Computer Science and Engineering for his immense care and encouragement towards us throughout the course of this project.

We also take this opportunity to express a deep sense of gratitude to our Supervisor **Dr.P.Vasuki** and our Project Co-ordinator **Dr.B.Selvapriya** for their cordial support, valuable information and guidance, They helped us in completing this project through various stages.

We thank our department faculty, supporting staff and friends for their help and guidance to complete this project.

B. Deepak (U20CS166)
B. Naga Sudhakar (U20CS156)
Ch. Rakesh (U20CS174)
B. Tirumalesh (U20CS214)

ABSTRACT

The three main factors that contribute to major driving errors are alcohol, sleepiness, and reckless driving. The Intelligent Transportation System's driver drowsiness detection system, which is the subject of this research, detects unusual driving behavior when the driver is using a computer. To ensure road safety, driving assistance systems must be able to assess a driver's level of alertness. Security. By watching blink patterns and eye movements, driver tiredness can be identified early enough to prevent crashes caused by drowsiness. The suggested solution uses computer vision techniques to provide a non-intrusive driver tiredness monitoring system. The simulation findings showed that even with the driver's glasses on and the car's interior lighting, the system was still able to identify fatigue. In addition, the gadget can identify drowsiness in less than two seconds. Alarms are used to correct the identified abnormal behavior in real time. Image analysis, image processing, and image pattern Edge detection has applications in human vision, computer vision, and recognition. In our project, a convolution neural network (CNN) technique is used to estimate the accuracy of drowsiness detection.

Keywords:-

CNN-Powered Solution, Reducing Accidents, Driver Alertness, Safety Technology

TABLE OF CONTENTS

Chapter No.	Title	Page No
	Bonafide Certificate	II
	Declaration	III
	Acknowledgments	IV
	Abstract	V
	Table Of Contents	VI
	List Of Figures	VIII
1	INTRODUCTION	1
1.1	Objective	1
1.2	Scope	2
1.3	Problem Identification	3
1.4	Methodology	4
2	LITERATURE SURVEY	5
2.1	Overview Of Key Research	5
2.2	Existing System	10
2.3	Requirement Specification	12
2.4	Feasibility Study	12
2.4.1	Economical Feasibility	12
2.4.2	Technical Feasibility	13
2.4.3	Social Feasibility	13
2.5	Innovation And Usefulness	13
3	PROPOSED METHODOLOGY	15
3.1	Proposed System	15
3.2	Software Process Model Used	18
3.3	Market Potential And Competitive Advantage	19
3.3.1	Distance And Color Directional Clustering Algorithm	19
3.3.2	Competitive Advantage	20
3.4	Challenges Of Stay Awake Stay Alive System	21
3.5	Machine Learning	22
3.5.1	Machine Learning Vs Traditional Programming	23
3.5.2	Machine Learning Algorithms and Where They are Used?	25
3.5.3	How To Choose Machine Learning Algorithm	27
3.5.4	Challenges And Limitations Of Machine Learning	28

3.5.5	Application Of Machine Learning	28
3.5.6	Example of Application of Machine Learning in supply Chain	29
3.5.7	Example Of Machine Learning Google Car	29
4	EXPERIMENTAL ANALYSIS AND DISCUSSION	33
4.1	Data Flow Diagram	34
4.2	Use case Diagram	35
4.3	Sequence Diagram	36
4.4	Result and Discussion	38
4.5	Requirement Analysis	40
4.6	System Design And Testing Plan	42
4.6.1	Image Processing	42
5	TECHNOLOGY AND TOOLS	72
5.1	Python	72
5.2	History	72
5.3	Opencv	77
5.3.1	Opencv History	77
5.3.2	Applications	78
5.4	Programming Language	79
5.4.1	Os Support	79
6	CONCLUSION AND FUTURE ENHANCEMENT	81
6.1	Conclusion	81
6.2	Future Enhancement	83
7	APPENDIX	85
7.1	Requirements	85
7.2	Source Code	85
7.3	Sample Output	90
	REFERENCES	91
	PUBLICATIONS	93

LIST OF FIGURES

Fig.NO	TITLE	PAGE NO
3.1	Block Diagram	15
3.2	Traditional Programming	23
3.3	Machine Learning	24
3.4	Learning Phase	24
3.5	Inference From Model	25
3.6	Machine Learning Algorithms	26
3.7	Algorithms	28
4.1	DFD Of Stay Awake Stay Alive	35
4.2	Use-Case Diagram	36
4.3	Sequence Diagram	37
4.4	Awake Result	40
4.5	Drowsiness Result	40
4.6	2×2 Pixel Area Displaying One Composite Color	47
4.7	Convolutional Neural Network	49
4.8	Small Image (Left) And Kernel (Right) To Illustrate Convolution	49
4.9	1st And 2nd Derivative Of Edge Illustrated In One Dimension	53
4.10	Convolution Kernel For Mean Filter With 3×3 Neighborhood.	57
4.11	Truth-Tables For And And Nand	58
4.12	Graylevel Image And The Corresponding Surface In Image Space	61
4.13	Two Connected Components Based On 4-Connectivity	65
4.14	Some Example Structuring Elements	68
4.15	Mapping Function For Wrapping The Pixel Values Of An 8-Bit Image	69
4.16	Mapping Function For Saturating An 8-Bit Image	70
7.1	Awake Result	90
7.2	Drowsiness Result	90

CHAPTER – 1

INTRODUCTION

In recent years, road accidents caused by driver drowsiness have become a critical concern worldwide. According to the World Health Organization (WHO), approximately 1.35 million people die each year as a result of road traffic crashes, with drowsy driving being a significant contributing factor. Identifying and preventing driver drowsiness is thus paramount for road safety and accident prevention.

In every country, road traffic accidents are a major public health problem and cause huge societal and financial burdens. Sleepiness causes disruption of neurological functions. Factors that contribute to the incidence of road traffic accidents range from continued driving even when feeling drowsy, having a physical condition, fewer sleeping hours, more working hours, and nutritional imbalances.

Several studies during the last 20 years have suggested that sleepiness is among the main factors that cause road traffic accidents. Sleepiness while driving contributes to 3% to > 30% of all road traffic accidents globally, which may involve a variety of sleep conditions but also may be caused by sleep deprivation.

The primary objective of this research is to design, implement, and evaluate a CNN-based drowsiness detection system integrated into existing vehicle safety mechanisms. By leveraging the power of deep learning and computer vision, we aim to enhance the effectiveness of drowsiness detection, ultimately contributing to a safer driving environment for all road users.

1.1 Objective

The objective of "CNN BASED STAY AWAKE STAY ALIVE SYSTEM FOR ACCIDENT PREVENTION", is to effectively reduce drowsy driving accidents. By leveraging CNN technology, the solution aims to detect signs of driver drowsiness in real-time and provide timely alerts or interventions to keep the driver awake and alert, thereby preventing potential accidents caused by fatigue-induced impairment.

The primary objective of our project is to employing Convolutional Neural Networks (CNN) for driver drowsiness detection is to enhance road safety by accurately identifying signs of driver fatigue or drowsiness in real-time. Specifically, the objectives include:

- **Accurate Detection:** Develop a CNN model capable of accurately detecting various indicators of drowsiness, such as drooping eyelids, yawning, and changes in facial expressions, with a high degree of precision and reliability.
- **Real-time Monitoring:** Implement the CNN-based system to continuously monitor the driver's facial features and expressions in real-time, enabling timely detection of drowsiness and immediate intervention to prevent potential accidents.
- **Low False Positive Rate:** Minimize false alarms by optimizing the CNN model to accurately distinguish between genuine instances of drowsiness and benign facial movements or temporary distractions, thereby maintaining the system's credibility and usability.
- **User-friendly Interface:** Develop a user-friendly interface for the drowsiness detection system, ensuring ease of deployment and interaction for both drivers and fleet operators, with clear visual or auditory alerts to prompt appropriate responses in case of detected drowsiness.

By achieving these objectives, the CNN-based drowsiness detection system aims to significantly reduce the incidence of road accidents caused by drowsy driving, thereby enhancing overall road safety and saving lives.

1.2 Scope

- **Real-time Monitoring:** CNN-based drowsiness detection systems can be implemented to continuously monitor drivers in real-time, allowing for immediate detection and intervention when signs of drowsiness are detected. This real-time capability enhances the system's effectiveness in preventing accidents by providing timely alerts to the driver.
- **Personalized Alerting Mechanisms:** The scope exists for developing personalized alerting mechanisms based on individual driver characteristics and behavior patterns. By leveraging CNN for drowsiness detection, systems can be tailored to account for differences in facial features, expressions, and driving styles, thereby improving the accuracy of alerting mechanisms.
- **Adaptability to Environmental Conditions:** CNN-based models can be trained to adapt to different environmental conditions, such as varying lighting conditions or changes in weather. This adaptability enhances the robustness of the drowsiness detection system, ensuring reliable performance across different driving scenarios.

- **Data Privacy and Security:** Given the sensitive nature of facial data used for drowsiness detection, there is a scope for research in ensuring data privacy and security. Techniques such as federated learning or edge computing can be explored to minimize the transmission of sensitive data and protect user privacy while still enabling effective drowsiness detection.
- **Validation and Certification:** Further research is needed to validate the performance of CNN-based drowsiness detection systems under real-world driving conditions and obtain certifications for deployment in commercial vehicles. This involves conducting large-scale trials and ensuring compliance with regulatory standards for automotive safety systems.

Overall, the scope for driver drowsiness detection using CNN is vast, with opportunities for innovation in technology, integration with existing systems, and addressing practical challenges such as data privacy and validation. By leveraging the capabilities of CNN, we can develop more reliable and effective solutions for preventing accidents caused by drowsy driving, ultimately saving lives and improving road safety.

1.3 Problem Identification

Identifying drowsiness in drivers is a critical challenge for road safety, as it poses a significant risk of accidents and fatalities. Traditional methods of detecting drowsiness, such as monitoring eyelid movements or steering patterns, have limitations in terms of accuracy and reliability, particularly in real-time scenarios. Furthermore, existing driver assistance systems often lack the capability to effectively discern subtle signs of drowsiness, leading to delayed or inadequate interventions. This problem is exacerbated by the increasing prevalence of distractions, such as smartphones and in-vehicle entertainment systems, which further contribute to driver fatigue and reduced attentiveness. Thus, there is a pressing need for more advanced and robust drowsiness detection mechanisms that can accurately identify signs of fatigue and alert drivers promptly to mitigate the risk of accidents. Addressing this problem requires innovative approaches that leverage cutting-edge technologies such as Convolutional Neural Networks (CNN) to analyze facial expressions and other visual cues indicative of drowsiness with high precision and efficiency. By addressing this problem, we can significantly enhance road safety and prevent countless accidents caused by drowsy driving.

1.4 Methodology

- **Data Collection:** Gather a diverse dataset of images or video clips featuring drivers in various states of alertness and drowsiness. Ensure the dataset includes sufficient variations in lighting conditions, facial expressions, and head poses. Annotate the dataset to label instances of drowsiness and non-drowsiness.
- **Data Preprocessing:** Resize all images to a uniform size suitable for CNN input. Normalize pixel values to a common scale. Augment the data set to increase its diversity and robustness, including techniques such as rotation, flipping, and adding noise.
- **CNN Architecture Design:** Design a CNN architecture suitable for image classification tasks, considering factors such as depth, filter sizes, and pooling strategies. Leverage pre-trained CNN models as a starting point to benefit from transfer learning. Modify the last few layers of the pre-trained model to adapt it to the specific task of drowsiness detection.
- **Training Procedure:** Split the data set into training, validation, and test sets, maintaining a balanced distribution of drowsy and non-drowsy samples in each set. Utilize techniques such as batch normalization and dropout regularization to improve generalization and prevent over fitting. Train the CNN model using a suitable optimization algorithm (e.g., stochastic gradient descent or Adam) with an appropriate learning rate. Monitor the training process using metrics such as loss and accuracy on the validation set and employ early stopping if performance plateaus.
- **Real-time Deployment:** Integrate the trained CNN model into a real-time drowsiness detection system, leveraging onboard cameras or other sensors installed in vehicles. Implement efficient algorithms for image preprocessing and inference to ensure low latency and real-time responsiveness. Validate the system's performance through extensive field testing under diverse driving conditions and scenarios.

CHAPTER - 2

LITERATURE SURVEY

Literature survey is the most important step in software development process. Before developing the tool it is necessary to determine the time factor, economy and company strength. Once these things are satisfied, then the next step is to determine which operating system and language can be used for developing the tool. Once the programmers start building the tool the programmers need lot of external support. This support can be obtained from senior programmers, from book or from websites. Before building the system the above consideration are taken into account for developing the proposed system. The major part of the project development sector considers and fully survey all the required needs for developing the project. For every project Literature survey is the most important sector in software development process. Before developing the tools and the associated designing it is necessary to determine and survey the time factor, resource requirement, man power, economy, and company strength. Once these things are satisfied and fully surveyed, then the next step is to determine about the software specifications in the respective system such as what type of operating system the project would require, and what are all the necessary software are needed to proceed with the next step such as developing the tools, and the associated operations.

2.1 Overview of key research

[1] Driver Drowsiness Prediction Based on Multiple Aspects Using Image Processing Techniques

Authors: VU Maheswari, R Aluvalu, MVVP Kantipudi, KK Chennam, K Kotecha, JR Saini

The majority of the accidents were happening perpetually due to driver drowsiness over the decades. Automation has been playing key role in many fields to provide conformity and improve the quality of life of the users. Though various drowsiness detection systems have been developed during last decade based on many factors, still the systems were demanding an improvement in terms of efficiency, accuracy, cost, speed, and availability, etc. In this paper, proposed an integrated approach depends on the Eye and mouth closure status (PERCLOS) along with the calculation of the new proposed vector FAR (Facial Aspect Ratio) similarly to EAR and MAR. This helps to find the status of the closed eyes or opened mouth like yawning, and any frame finds that has hand gestures like nodding or covering opened mouth with hand as innate nature of humans when trying to control the sleepiness.

[2] An Intelligent Driving Assistance System Based on Lightweight Deep Learning models

Authors: KF Lee, XZ Chen, CW Yu, KY Chin, YC Wang, CY Hsiao, YL Chen

An intelligent driver assistance system is developed in this study, which is able to remind the drivers to turn on the head lights or wipers through situation recognition method when driving at night or on rainy days. Furthermore, the object detection results from multiple perspective views are integrated, and the surrounding object detection results are produced for collision avoidance. The system is able to alarm the drivers based on the lightweight deep learning model and the distance estimation method when surrounding vehicles are too close. Experimental results show that the proposed methods and the chosen lightweight model in our proposed system obtain reliable performance and sufficient computational efficiency under limited computing resource. In conclude, our proposed system obtains high probability to be adopted for the development of advanced driver assistance systems (ADAS). The proposed system can not only assist the driver in determining the vision ahead, but also provide an instant overview of the vehicle's surrounding conditions to enhance driving safety.

[3] Early Identification and Detection of Driver Drowsiness by Hybrid Machine Learning

Authors: A Altameem, A Kumar, RC Poonia, S Kumar, AKJ Saudagar

Drunkenness or exhaustion is a leading cause of car accidents, with severe implications for road safety. Several drowsiness detection technologies to monitor for signs of inattention while driving and notifying the driver can be adopted. Sensors in self-driving cars must detect if a driver is sleepy, angry, or experiencing extreme changes in their emotions, such as anger. These sensors must constantly monitor the driver's facial expressions and detect facial landmarks in order to extract the driver's state of expression presentation and determine whether they are driving safely

[4] Real-Time Driver-Drowsiness Detection System Using Facial Features

Authors: W Deng, R Wu

The face, an important part of the body, conveys a lot of information. When a driver is in a state of fatigue, the facial expressions, e.g., the frequency of blinking and yawning, are different from those in the normal state. In this paper, we propose a system called DriCare,

which detects the drivers' fatigue status, such as yawning, blinking, and duration of eye closure, using video images, without equipping their bodies with devices. Owing to the shortcomings of previous algorithms, we introduce a new face-tracking algorithm to improve the tracking accuracy. Further, we designed a new detection method for facial regions based on 68 key points. Then we use these facial regions to evaluate the drivers' state. By combining the features of the eyes and mouth, DriCare can alert the driver using a fatigue warning. The experimental results showed that DriCare achieved around 92% accuracy.

[5] Driver Drowsiness Detection Using Deep Learning

Authors: Yeresime Suresh; Rashi Khandelwal; Matam Nikitha; Mohammed Fayaz; Vinaya Soudhri

Driver drowsiness is a major cause of road accidents worldwide, and a drowsy driver is a serious threat to road safety. To address this issue, researchers have developed various driver drowsiness detection systems that can alert the driver before a mishap occurs. In this paper, we present a comprehensive review of driver drowsiness detection systems, including their underlying techniques, advantages, and limitations. We examine the most commonly used techniques for detecting driver drowsiness, such as physiological measures, eye-tracking, and machine learning approaches. We also highlight the challenges associated with the development and implementation of such systems, including variability in individual sleep patterns, changing environmental conditions, and the trade-off between accuracy and user-friendliness. This review aims to provide a critical analysis of the state-of-the-art in driver drowsiness detection systems, and to identify research gaps and future directions for improving road safety.

[6] Driver Drowsiness Detection Using Machine Learning Algorithm

Authors: N Prasath, J Sreemathy, P Vigneshwaran

It is just a short time until self-driving vehicles become omnipresent; be that as it may, human driving management will stay a need for quite a long time. Driver Drowsiness is one of the significant reasons of roadways accidents these days. Hence fatigue and drowsiness detection play a major role in preventing the road accidents. Every year, because of this there is an increase in the number of deaths and injuries globally. Recently, in this decade, many images processing-based approaches were created and used to detect driver's drowsiness status. To minimize the number of accidents, a method is proposed in this paper. The algorithm focuses on the eye closure and yawning ratios. The driver is alarmed, if he/she is feeling sleepy.

[7] Drivers' Drowsiness Detection and Warning Systems for Critical Infrastructures

Authors: IR Adochiei, OI Știrbu, NI Adochiei, M Pericle-Gabriel, CM Larco, SM Mustata, D Costin

Road traffic accidents, due to driver fatigue, tend to inflict high mortality rates comparing with accidents involving rested drivers. Currently there is an emerging automotive industry trend towards equipping vehicles with various driver-assistance technologies. Third parties also started producing complementary systems, including ones that can detect the driver's degree of fatigue, but this growing field requires further research and development. The main purpose of this paper is the development and implementation of a system capable to detecting and alert, in real-time, the driver's level of fatigue. Hence, in this work the authors are focused on the video monitoring of the driver face, especially on his eyes position in time, when open or closed, using a machine learning object detection algorithm, the Haar Cascade.

[8] Drivers Drowsiness Detection using Image Processing and I-Ear Techniques

Authors: S Ananthi, R Sathya, K Vaidehi, G Vijaya

Fatigue and micro-sleep at the wheel are often the cause of serious accidents. Consequently, the initial signs of micro-sleep can be detected before a critical situation arises. Drowsiness detection technique helps to prevent accidents caused by the drowsy drivers. Driver drowsiness detection involves three methodologies namely face detection, eye detection, drowsiness detection. Face detection technique is used to detect the driver's face by extracting the facial features using Haar Cascade Classifier Algorithm also it detects the person's eye region. This paper proposed to Identifying Eye Aspect Ratio (I-EAR) method to identify the eye closure of vehicle's driver. If the driving person's eye is closed for minimum of 40 frames or eye blink ratio is less than 8 per minute, then the proposed system identifies that the vehicle driving person is drowsy. Consequently, it began to alert sound intimation that vehicle driver is drowsy. This alert may help the drivers as a wake up call to take an immediate action by taking rest/refreshment and continue to drive safely.

[9] EEG-based System Using Deep Learning and Attention Mechanism for Driver Drowsiness Detection

Authors: Miankuan Zhu; Haobo Li; Jiangfan Chen; Mitsuhiro Kamezaki; Zutao Zhang; Zexi Hua; Shigeki Sugano

The lack of sleep (typically <6 hours a night) or driving for a long time are the reasons of drowsiness driving and caused serious traffic accidents. With pandemic of the

COVID-19, drivers are wearing masks to prevent infection from it, which makes visual-based drowsiness detection methods difficult. This paper presents an EEG-based driver drowsiness estimation method using deep learning and attention mechanism. First of all, an 8-channels EEG collection hat is used to acquire the EEG signals in the simulation scenario of drowsiness driving and normal driving. Then the EEG signals are pre-processed by using the linear filter and wavelet threshold denoising. Secondly, the neural network based on attention mechanism and deep residual network (ResNet) is trained to classify the EEG signals. Finally, an early warning module is designed to sound an alarm if the driver is judged as drowsy. The system was tested under simulated driving environment and the drowsiness detection accuracy of the test set was 93.35%. Drowsiness warning simulation also verified the effectiveness of proposed early warning module.

[10] Robust Two-Stream Multi-Features Network for Driver Drowsiness Detection

Authors: Qi Shen, Shengjie Zhao, Rongqing Zhang, Bin Zhang

Drowsiness driving is a major cause of traffic accidents, and thus numerous previous researches have focused on driver drowsiness detection. Many drive relevant factors have been taken into consideration for fatigue detection. They can lead to high precision, but there are still several serious constraints, such as most existing models are environmentally susceptible. In this paper, fatigue detection is considered as a temporal action detection problem instead of image classification. The proposed detection system can be divided into four parts: (1) Localize the key patches of the detected driver facial picture for fatigue detection and calculate the corresponding optical flow. (2) Contrast Limited Adaptive Histogram Equalization (CLAHE) is used in our system to reduce the impact of different light conditions. (3) Three individual two-stream networks combined with 3D attention mechanism and 3D depthwise separable convolution are designed for each feature to extract temporal information. (4) The outputs of the three sub-networks will be concatenated and sent to the fully-connected network, which judges the status of the driver. The drowsiness detection system is trained and evaluated on the famous National Tsinghua University Driver Drowsiness Detection (NTHU-DDD) dataset and we obtain an accuracy of 94.46%, which outperforms most existing fatigue detection models.

[11] Driver stress detection via multimodal fusion using attention-based CNN-LSTM

Author: Chao Zhou, Pengfei Zhao, Bahareh Nakisa, Mohammad Naim Rastgoo, Ramesh Jain, Wen Gao

Stress has been identified as one of major contributing factors in car crashes due to its negative impact on driving performance. It is in urgent need that the stress levels of drivers can be detected in real time with high accuracy so that intervening or navigating measures can be taken in time to mitigate the situation. Existing driver stress detection models mainly rely on traditional machine learning techniques to fuse multimodal data. However, due to the non-linear correlations among modalities, it is still challenging for traditional multimodal fusion methods to handle the real-time influx of complex multimodal and high dimensional data, and report drivers' stress levels accurately. To solve this issue, a framework of driver stress detection through multimodal fusion using attention based deep learning techniques is proposed in this paper. Specifically, an attention based convolutional neural networks (CNN) and long short-term memory (LSTM) model is proposed to fuse non-invasive data, including eye data, vehicle data, and environmental data. Then, the proposed model can automatically extract features separately from each modality and give different levels of attention to features from different modalities through self-attention mechanism. To verify the validity of the proposed method, extensive experiments have been carried out on our dataset collected using an advanced driving simulator. Experimental results demonstrate that the performance of the proposed method on driver stress detection outperforms the state-of-the-art models with an average accuracy of 95.5%.

2.2 Existing System

1) Drowsy Driver Detection System Using Convolutional Neural Networks

This system employs CNNs to analyze facial expressions and eye movements to detect drowsiness in drivers.

2) Real-Time Driver Drowsiness Detection System

This system uses CNNs to process images from a camera pointed at the driver's face to detect signs of drowsiness such as eye closure and head nods.

3) Driver Fatigue Detection System Based on Deep Convolutional Neural Network

This system utilizes CNNs to extract features from facial images and classify the driver's state as either alert or drowsy.

4) Real-Time Driver Drowsiness Detection Using CNN

This system employs a CNN architecture to analyze facial images and detect

signs of driver drowsiness in real-time.

5) Driver Drowsiness Detection System Based on Convolutional Neural Networks and Transfer Learning

This system utilizes CNNs with transfer learning to detect drowsiness by analyzing facial features and expressions.

6) CNN-Based Real-Time Drowsiness Detection System for Drivers

This system uses CNNs to analyze facial images and classify the driver's state as either alert or drowsy in real-time.

7) Driver Drowsiness Detection Using CNN and LSTM

This system combines CNNs with Long Short-Term Memory (LSTM) networks to analyze both spatial and temporal features from facial images for drowsiness detection.

8) Driver Drowsiness Detection System Using CNN and Eye Gaze Tracking

This system integrates CNNs with eye gaze tracking to detect drowsiness based on both facial expressions and eye movements.

9) Driver Drowsiness Detection Based on Multi-Modal CNN

This system incorporates multiple CNNs trained on different modalities such as facial images and infrared images to improve drowsiness detection accuracy.

10) Real-Time Driver Drowsiness Detection Using Deep CNN and Facial Landmarks

This system employs deep CNNs along with facial landmark detection to analyze facial expressions and detect drowsiness in real-time.

11) Driver Drowsiness Detection System Based on Multi-Task CNN

This system utilizes a multi-task CNN architecture to simultaneously detect facial landmarks and classify the driver's drowsiness state.

12) Driver Fatigue Detection Using CNN and Eye State Analysis

This system combines CNNs with analysis of eye states (open or closed) to detect driver fatigue and prevent accidents.

13) Real-Time Driver Drowsiness Detection Using CNN and Optical Flow

This system integrates CNNs with optical flow analysis to capture both spatial and temporal features from facial images for drowsiness detection.

14) Driver Drowsiness Detection System Based on Attention Mechanism in CNN

This system incorporates attention mechanisms into CNNs to dynamically weight facial features for more accurate drowsiness detection.

15) Driver Drowsiness Detection System Using Hybrid CNN-RNN Architecture

This system employs a hybrid architecture combining CNNs for feature extraction and Recurrent Neural Networks (RNNs) for temporal sequence analysis to detect driver drowsiness.

2.3 Requirement Specification

■ HARDWARE REQUIREMENTS

- System : Pentium i3 Processor
- Hard Disk : 500 GB.
- Monitor : 15'' LED
- Input Devices : Keyboard, Mouse
- Ram : 2 GB

■ SOFTWARE REQUIREMENTS

- Operating system : Windows 10
- Coding Language :python

2.4 Feasibility Study

The feasibility of the project is analyzed in this phase and business proposal is put forth with a very general plan for the project and some cost estimates. During system analysis the feasibility study of the proposed system is to be carried out. This is to ensure that the proposed system is not a burden to the company. For feasibility analysis, some understanding of the major requirements for the system is essential.

Three key considerations involved in the feasibility analysis are

- ECONOMICAL FEASIBILITY
- TECHNICAL FEASIBILITY
- SOCIAL FEASIBILITY

2.4.1 ECONOMICAL FEASIBILITY

This study is carried out to check the economic impact that the system will have on the organization. The amount of fund that the company can pour into the research and development of the system is limited. The expenditures must be justified. Thus the developed system as well within the budget and this was achieved because most of the technologies used are freely available. Only the customized products had to be purchased.

2.4.2 TECHNICAL FEASIBILITY

This study is carried out to check the technical feasibility, that is, the technical requirements of the system. Any system developed must not have a high demand on the

available technical resources. This will lead to high demands on the available technical resources. This will lead to high demands being placed on the client. The developed system must have a modest requirement, as only minimal or null changes are required for implementing this system.

2.4.3 SOCIAL FEASIBILITY

The aspect of study is to check the level of acceptance of the system by the user. This includes the process of training the user to use the system efficiently. The user must not feel threatened by the system, instead must accept it as a necessity. The level of acceptance by the users solely depends on the methods that are employed to educate the user about the system and to make him familiar with it. His level of confidence must be raised so that he is also able to make some constructive criticism, which is welcomed, as he is the final user of the system.

2.5 INNOVATION AND USEFULNESS

The integration of Convolutional Neural Networks (CNN) for driver drowsiness detection represents a groundbreaking innovation with profound implications for road safety and accident prevention. Unlike traditional methods which rely on simple metrics like eyelid movements or steering patterns, CNN-based systems offer a sophisticated, real-time analysis of facial features and expressions, providing a more accurate and reliable means of identifying drowsiness in drivers.

The key innovation lies in the utilization of deep learning techniques to interpret complex visual data captured by on board cameras. CNN are particularly adept at processing image-based information, making them ideal for tasks like facial recognition and expression analysis. By leveraging the power of CNN, drowsiness detection systems can effectively discern subtle cues indicative of driver fatigue, such as drooping eyelids, yawning, and changes in facial expression. This enables timely interventions, such as alerts or automated vehicle control adjustments, to prevent accidents before they occur.

The usefulness of CNN-based drowsiness detection extends far beyond its technological sophistication. It addresses a critical safety concern on our roads, where drowsy driving contributes significantly to accidents and fatalities. By providing early warnings to drivers, these systems help mitigate the risks associated with fatigue-related impairment, potentially saving countless lives and reducing the economic burden of road accidents.

Moreover, CNN-based drowsiness detection systems offer versatility and scalability, making them applicable across various vehicle types and driving conditions. Whether in

commercial trucks, public transportation, or personal vehicles, these systems can adapt to different environments and driver behaviors, enhancing safety across diverse transportation settings.

Furthermore, the non-intrusive nature of CNN-based drowsiness detection is a significant advantage. Unlike some existing solutions that require cumbersome sensors or physical contact with the driver, CNN-based systems rely solely on on board cameras, minimizing disruptions to the driving experience while still providing effective monitoring capabilities. This makes them more user-friendly and widely acceptable among drivers and vehicle manufacturers.

Additionally, the continuous learning capability of CNN contributes to the ongoing improvement of drowsiness detection algorithms. As more data is collected and analyzed, CNN-based systems can refine their understanding of drowsiness indicators, further enhancing their accuracy and reliability over time. This adaptability ensures that the technology remains at the forefront of driver safety, even as driving habits and environmental factors evolve.

In conclusion, the integration of Convolutional Neural Networks for driver drowsiness detection represents a significant leap forward in road safety technology. By harnessing the power of deep learning and computer vision, CNN-based systems offer unparalleled accuracy, reliability, and versatility in identifying and preventing drowsy driving accidents. With the potential to save lives, reduce injuries, and enhance overall transportation .

CHAPTER - 3

PROPOSED METHODOLOGY

In recent years, the automotive industry has witnessed a significant surge in the integration of advanced driver assistance systems (ADAS) to enhance safety and mitigate potential risks associated with driver fatigue and drowsiness. Among these systems, the development of driver drowsiness detection systems has garnered substantial attention due to their potential to prevent accidents caused by impaired driving. Convolutional Neural Networks (CNN) have emerged as a powerful tool in computer vision tasks, demonstrating remarkable capabilities in image recognition and classification. Leveraging the inherent capacity of CNN, this proposed methodology aims to devise an efficient and reliable system for detecting driver drowsiness based on visual cues extracted from in-car monitoring systems, such as dashboard cameras or infrared sensors.

3.1 Proposed System

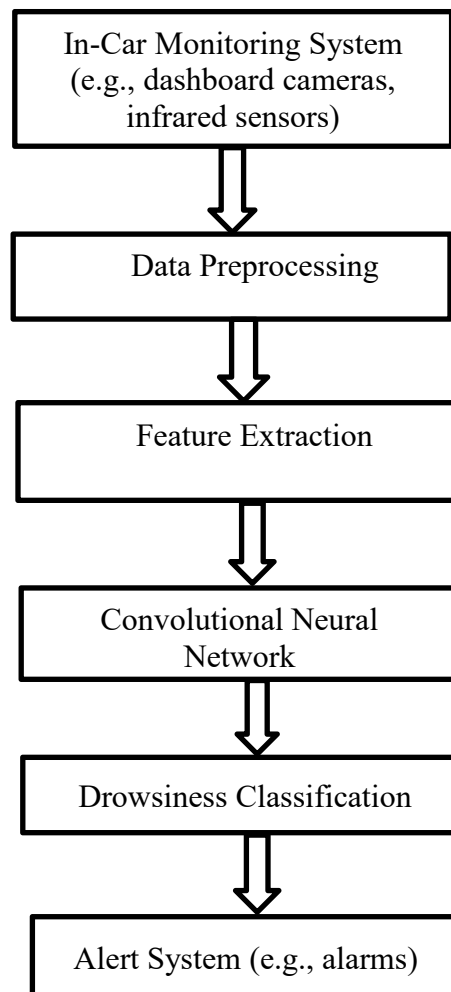


Figure 3.1: Block Diagram

Data Collection

Gather a dataset of images or videos containing both alert and drowsy states of drivers. These images or videos should ideally cover various lighting conditions, different facial expressions, and diverse individuals to ensure the model's robustness.

Data Preprocessing

Preprocess the collected data to enhance the model's performance. This may involve tasks such as resizing images to a consistent resolution, normalization, and augmentation techniques like rotation, flipping, or adjusting brightness to increase the diversity of the training data.

Model Architecture Selection

Choose an appropriate CNN architecture for the drowsiness detection task. Common choices include modified versions of well-known architectures like VGG, ResNet, or Inception. Alternatively, you can design a custom architecture tailored to the specific requirements of the task.

Model Training

Split the preprocessed dataset into training, validation, and testing sets. Train the CNN model using the training set and validate its performance using the validation set. This step involves adjusting hyperparameters, such as learning rate, batch size, and number of epochs, to optimize the model's performance.

Model Evaluation

Evaluate the trained model's performance using the testing set. Metrics such as accuracy, precision, recall, and F1-score can be used to assess the model's effectiveness in detecting drowsiness.

Data Acquisition and Preprocessing

Acquire a diverse dataset consisting of video recordings capturing driver behavior under various driving conditions. Preprocess the data to extract relevant frames, align facial landmarks, and normalize images to ensure consistency and enhance model performance.

Feature Extraction

Utilize pre-trained CNN architectures (e.g., VGG16, ResNet) as feature extractors to capture hierarchical features from facial images. Employ techniques such as transfer learning to fine-tune the network on the specific task of drowsiness detection, adapting to the nuances of driver behavior.

Model Development

Design and train a CNN-based classification model to discern between alert and drowsy states based on extracted features. Explore various network architectures, including multi-stream CNN for integrating temporal information from sequential frames.

Model Evaluation and Validation

Assess the performance of the trained model using appropriate evaluation metrics such as accuracy, precision, recall, and F1-score. Conduct cross-validation and potentially deploy techniques like data augmentation to ensure robustness and generalization.

Real-time Integration

Implement the trained model within an embedded system capable of processing streaming video input in real-time. Integrate the drowsiness detection system with existing ADAS frameworks or standalone in-car monitoring systems.

Deployment

Once the model achieves satisfactory performance, integrate it into a real-time system for driver drowsiness detection. This system can be implemented using hardware such as cameras installed in vehicles to capture real-time images of drivers' faces. The deployed model should be capable of processing these images efficiently and providing timely alerts when drowsiness is detected.

Alert Mechanism

Implement a mechanism to alert the driver when drowsiness is detected. This can be done through auditory or visual cues, such as alarms, vibrations, or messages displayed on the vehicle's dashboard.

Safety Considerations

Ensure that the drowsiness detection system does not distract the driver or interfere with their ability to safely operate the vehicle. Implement fail-safe mechanisms to prevent false alarms or erroneous interventions that could potentially endanger the driver or other road users.

Regulatory Compliance

Ensure that the system complies with relevant regulations and standards for automotive safety and driver assistance systems in the target region or market. Conduct thorough testing and validation to demonstrate the effectiveness and reliability of the system in real-world conditions.

Continuous Improvement

Continuously monitor the system's performance in real-world scenarios and collect feedback to identify areas for improvement. This may involve retraining the model with additional data or fine-tuning its parameters to enhance its accuracy and robustness.

3.2 Software Process Model Used

Iterative Model

The Iterative model allows for repetitive cycles of development and refinement. It's suitable for projects where requirements may evolve over time or where there's a need for continuous improvement. In this model, you could start with basic CNN models for drowsiness detection, test them iteratively with real-world data, gather feedback, and refine the models accordingly.

Agile Model

Agile methodologies like Scrum or Kanban could be beneficial for projects where requirements are not completely known upfront or may change during development. Agile allows for frequent collaboration between developers and stakeholders, enabling quick adaptations to changing needs. It's suitable for projects where rapid prototyping and incremental development are crucial.

Spiral Model

The Spiral model combines elements of both iterative and waterfall models. It's suitable for projects with high-risk factors, where continuous refinement is necessary. Each iteration in the Spiral model involves identifying risks, developing prototypes, and testing them. This model could be beneficial for a drowsiness detection system, where safety is a significant concern, and iterative improvements are essential.

V-Model

The V-Model is a variation of the waterfall model, emphasizing the verification and validation of each stage of development. It's suitable for projects with well-defined requirements upfront and a focus on rigorous testing. In the context of a drowsiness detection system, the V-Model ensures that each stage of CNN model development, from requirements specification to testing and deployment, is thoroughly verified and validated.

DevOps Model

DevOps focuses on the integration of development and operations teams, emphasizing automation, continuous integration, and continuous delivery. For a drowsiness detection system based on CNN, the DevOps model ensures efficient collaboration between developers and operations teams, facilitating rapid deployment and updates.

3.3 Market Potential and Competitive Advantage

3.3.1 Market Potential

Automotive Industry

There's a growing demand for advanced driver assistance systems (ADAS) to enhance vehicle safety. Regulations mandating safety features like drowsiness detection are increasing, particularly in developed regions. With the rise of autonomous vehicles, ensuring driver alertness becomes even more critical during the transition phase.

Transportation Sector

Companies operating fleets of vehicles (trucking, logistics, ridesharing) are keen on technologies that ensure driver attentiveness to reduce accidents and improve efficiency. Public transportation agencies could integrate drowsiness detection systems to enhance passenger safety.

Consumer Electronics

Integration of drowsiness detection into consumer vehicles and wearables can attract safety-conscious consumers. It could be part of a broader suite of features in smart vehicles.

Healthcare

Drowsiness detection systems can also be integrated into medical devices aimed at monitoring the health of individuals prone to conditions affecting alertness (e.g., sleep disorders, certain medications).

3.3.2Competitive Advantage

Accuracy and Reliability

CNNs excel in feature extraction from visual data, enabling highly accurate drowsiness detection. Continuous learning capabilities of CNNs allow for refinement and adaptation to diverse driving conditions and individual differences.

Real-time Processing

CNN-based systems can process video input in real-time, allowing for immediate alerts when signs of drowsiness are detected. This real-time response is crucial for preventing accidents caused by driver fatigue.

Scalability

CNN-based solutions can be scaled to accommodate various types of vehicles, from personal cars to commercial trucks. The adaptability of CNN architectures allows for customization based on different vehicle types and driver preferences.

Integration Potential

Integration with existing ADAS or vehicle safety systems can leverage CNN-based drowsiness detection as part of a comprehensive safety suite. Compatibility with other sensor data (e.g., steering wheel movements, physiological sensors) can enhance the accuracy of drowsiness detection.

Cost-effectiveness

As CNN-based technologies mature and hardware becomes more powerful, the cost of implementing drowsiness detection systems is likely to decrease. The potential cost savings from preventing accidents and reducing insurance premiums make this technology highly attractive to both manufacturers and end-users.

3.4 Challenges of Stay Awake Stay Alive System

Detecting driver drowsiness using Convolutional Neural Networks (CNNs) presents both opportunities and challenges in enhancing road safety. CNNs have demonstrated remarkable capabilities in image recognition tasks, making them a promising tool for analyzing visual cues indicative of driver fatigue. However, several challenges persist in effectively implementing CNN-based drowsiness detection systems.

One of the primary challenges is the variability in facial expressions and head poses among drivers. CNNs rely on learning features from input images, but variations in facial orientation and lighting conditions can significantly affect the network's ability to accurately detect drowsiness. Preprocessing techniques such as normalization and data augmentation may help mitigate these challenges, but they may not completely eliminate the issue.

Another obstacle is the limited availability of diverse and annotated datasets for training CNN models. Annotated datasets containing a wide range of facial expressions and drowsiness indicators are crucial for training robust models. However, collecting and labeling such datasets can be time-consuming and resource-intensive. Moreover, biases in the data, such as an overrepresentation of certain demographic groups, can lead to algorithmic biases and affect the model's generalization capabilities.

Furthermore, real-time processing requirements impose constraints on the computational efficiency of CNN-based drowsiness detection systems. While CNNs offer impressive accuracy, their computational complexity can be prohibitive for deployment in resource-constrained environments such as embedded systems in vehicles. Optimizing model architectures and leveraging hardware accelerators like GPUs and TPUs can help address this challenge, but striking a balance between accuracy and computational efficiency remains a significant concern.

Additionally, ensuring the robustness and reliability of CNN-based drowsiness detection systems in real-world scenarios poses a considerable challenge. Factors such as

varying environmental conditions, occlusions, and artifacts in input data can affect the performance of the model. Adversarial attacks, where slight perturbations to input images can lead to incorrect predictions, are also a concern. Developing techniques for adversarial training and model robustification is essential to enhance the resilience of CNN-based drowsiness detection systems.

Moreover, addressing privacy concerns associated with facial data collection and processing is paramount. Deploying CNN-based drowsiness detection systems raises privacy implications, as it involves capturing and analyzing individuals' facial features in real-time. Implementing privacy-preserving techniques such as on-device processing, anonymization, and secure data transmission can help alleviate these concerns and foster user trust in the technology.

Despite these challenges, continued research and development efforts hold promise for advancing CNN-based drowsiness detection systems. Collaborative initiatives involving researchers, industry stakeholders, and regulatory bodies are essential for addressing technical, ethical, and regulatory challenges. By overcoming these hurdles, CNN-based drowsiness detection systems have the potential to significantly enhance road safety by alerting drivers in real-time and preventing accidents caused by drowsy driving.

3.5 Machine Learning

What is Machine Learning?

Machine Learning is a system of computer algorithms that can learn from example through self-improvement without being explicitly coded by a programmer. Machine learning is a part of artificial Intelligence which combines data with statistical tools to predict an output which can be used to make actionable insights.

The breakthrough comes with the idea that a machine can singularly learn from the data (i.e., example) to produce accurate results. Machine learning is closely related to data mining and Bayesian predictive modeling. The machine receives data as input and uses an algorithm to formulate answers.

A typical machine learning tasks are to provide a recommendation. For those who have a Netflix account, all recommendations of movies or series are based on the user's historical data. Tech companies are using unsupervised learning to improve the user experience with personalizing recommendation.

Machine learning is also used for a variety of tasks like fraud detection, predictive maintenance, portfolio optimization, automatize task and so on.

3.5.1 Machine Learning vs. Traditional Programming

Traditional programming differs significantly from machine learning. In traditional programming, a programmer code all the rules in consultation with an expert in the industry for which software is being developed. Each rule is based on a logical foundation; the machine will execute an output following the logical statement. When the system grows complex, more rules need to be written. It can quickly become unsustainable to maintain.

Traditional programming differs significantly from machine learning. In traditional programming, a programmer code all the rules in consultation with an expert in the industry for which software is being developed. Each rule is based on a logical foundation; the machine will execute an output following the logical statement. When the system grows complex, more rules need to be written. It can quickly become unsustainable to maintain.




Figure 3.2 : Traditional Programming

Machine learning is supposed to overcome this issue. The machine learns how the input and output data are correlated and it writes a rule. The programmers do not need to write new rules each time there is new data. The algorithms adapt in response to new data and experiences to improve efficacy over time.




Figure 3.3 : Machine Learning

How does Machine Learning Work?

Machine learning is the brain where all the learning takes place. The way the machine learns is similar to the human being. Humans learn from experience. The more we know, the more easily we can predict. By analogy, when we face an unknown situation, the likelihood of success is lower than the known situation. Machines are trained the same. To make an accurate prediction, the machine sees an example. When we give the machine a similar example, it can figure out the outcome. However, like a human, if its feed a previously unseen example, the machine has difficulties to predict.

The core objective of machine learning is the **learning** and **inference**. First of all, the machine learns through the discovery of patterns. This discovery is made thanks to the **data**. One crucial part of the data scientist is to choose carefully which data to provide to the machine. The list of attributes used to solve a problem is called a **feature vector**. You can think of a feature vector as a subset of data that is used to tackle a problem.

The machine uses some fancy algorithms to simplify the reality and transform this discovery into a **model**. Therefore, the learning stage is used to describe the data and summarize it into a model.




Figure 3.4 : Learning Phase

For instance, the machine is trying to understand the relationship between the wage of an individual and the likelihood to go to a fancy restaurant. It turns out the machine finds a positive relationship between wage and going to a high-end restaurant: This is the model

Inferring

When the model is built, it is possible to test how powerful it is on never-seen-before data. The new data are transformed into a features vector, go through the model and give a prediction. This is all the beautiful part of machine learning. There is no need to update the rules or train again the model. You can use the model previously trained to make inference on new data.




Figure 3.5 : Inference from Model

The life of Machine Learning programs is straightforward and can be summarized in the following points:

1. Define a question
2. Collect data
3. Visualize data
4. Train algorithm
5. Test the Algorithm
6. Collect feedback
7. Refine the algorithm
8. Loop 4-7 until the results are satisfying
9. Use the model to make a prediction

Once the algorithm gets good at drawing the right conclusions, it applies that knowledge to new sets of data.

3.5.2 Machine Learning Algorithms and Where they are Used?

Machine learning algorithms are utilized across various domains and industries to extract insights from data, make predictions, and automate decision-making processes. One prominent algorithm is the Support Vector Machine (SVM), widely used in classification tasks like spam email detection, medical diagnosis, and sentiment analysis. Decision Trees and Random Forests are employed in both classification and regression problems, such as credit risk assessment and stock market prediction. Neural Networks, including Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs), excel in

image recognition, natural language processing, and time series forecasting. K-Means clustering finds applications in customer segmentation, anomaly detection, and image compression. Gaussian Mixture Models (GMMs) are utilized in speech recognition, while Principal Component Analysis (PCA) aids in dimensionality reduction for data visualization and feature extraction. Reinforcement Learning algorithms like Q-Learning and Deep Q-Networks are employed in autonomous systems, gaming, and robotics. Additionally, Gradient Boosting Machines (GBMs) are popular in predictive modeling and recommendation systems. Overall, machine learning algorithms play a crucial role in enhancing efficiency, accuracy, and automation across a wide array of industries, from healthcare and finance to transportation and entertainment.

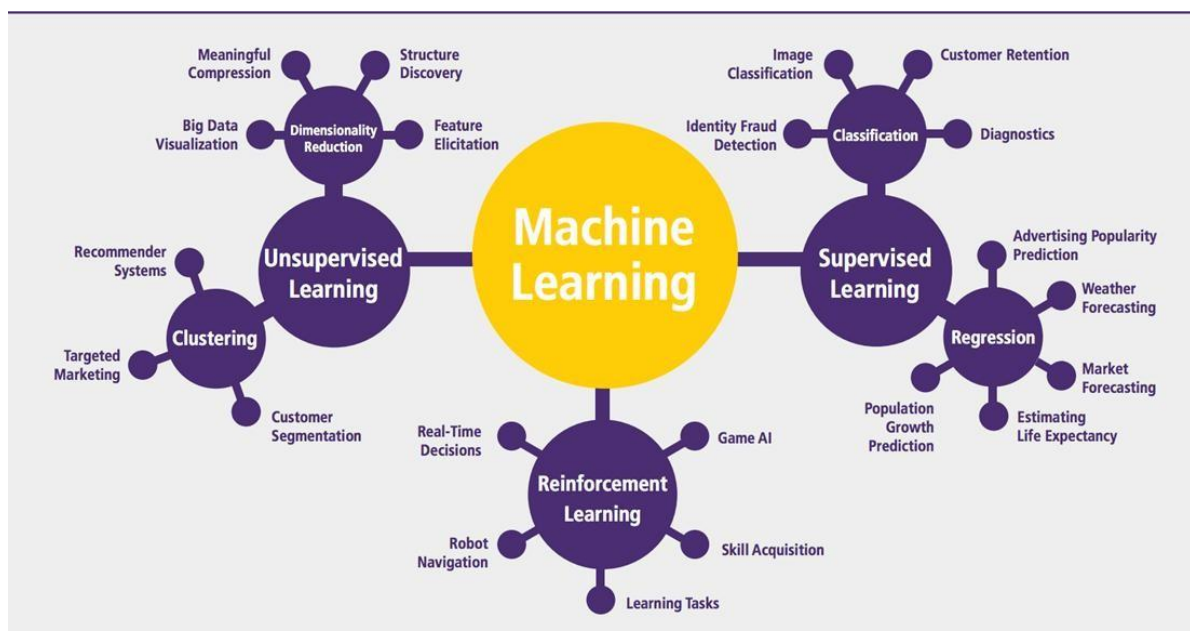


Figure 3.6 : Machine learning Algorithms

Machine learning can be grouped into two broad learning tasks: Supervised and Unsupervised. There are many other algorithms

Supervised learning

An algorithm uses training data and feedback from humans to learn the relationship of given inputs to a given output. For instance, a practitioner can use marketing expense and weather forecast as input data to predict the sales of cans.

You can use supervised learning when the output data is known. The algorithm will predict new data.

There are two categories of supervised learning:

- Classification task
- Regression task

Classification

Imagine you want to predict the gender of a customer for a commercial. You will start gathering data on the height, weight, job, salary, purchasing basket, etc. from your customer database. You know the gender of each of your customer, it can only be male or female. The objective of the classifier will be to assign a probability of being a male or a female (i.e., the label) based on the information (i.e., features you have collected). When the model learned how to recognize male or female, you can use new data to make a prediction. For instance, you just got new information from an unknown customer, and you want to know if it is a male or female. If the classifier predicts male = 70%, it means the algorithm is sure at 70% that this customer is a male, and 30% it is a female.

The label can be of two or more classes. The above Machine learning example has only two classes, but if a classifier needs to predict object, it has dozens of classes (e.g., glass, table, shoes, etc. each object represents a class)

Regression

When the output is a continuous value, the task is a regression. For instance, a financial analyst may need to forecast the value of a stock based on a range of feature like equity, previous stock performances, macroeconomics index. The system will be trained to estimate the price of the stocks with the lowest possible error.

Unsupervised learning

In unsupervised learning, an algorithm explores input data without being given an explicit output variable (e.g., explores customer demographic data to identify patterns) You can use it when you do not know how to classify the data, and you want the algorithm to find patterns and classify the data for you

3.5.3 How to Choose Machine Learning

AlgorithmMachine Learning (ML) algorithm

There are plenty of machine learning algorithms. The choice of the algorithm is based on the objective. In the Machine learning example below, the task is to predict the type of flower among the three varieties. The predictions are based on the length and the width of the petal. The picture depicts the results of ten different algorithms. The picture on the top left is the dataset. The data is classified into three categories: red, light blue and dark blue. There are some groupings. For instance, from the second image, everything in the upper left belongs to the red category, in the middle part, there is a mixture of uncertainty and light blue while the bottom corresponds to the dark category. The other images show different algorithms and how they try to classified the data.

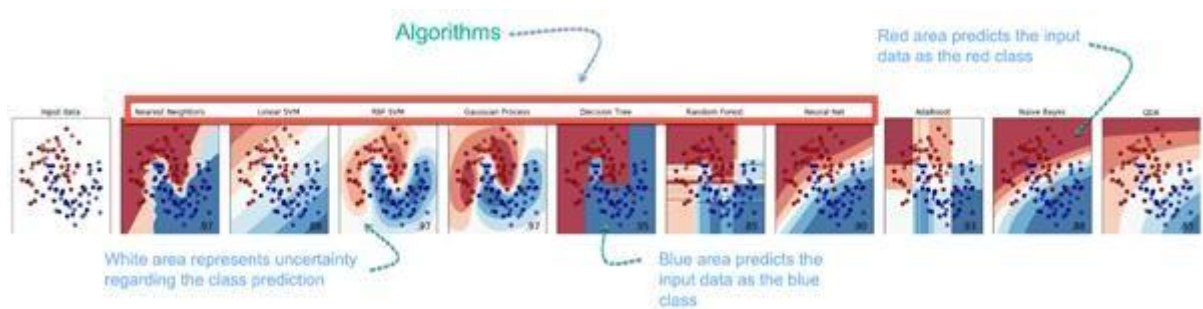


Figure 3.7 : Algorithms

3.5.4 Challenges and Limitations of Machine Learning

The primary challenge of machine learning is the lack of data or the diversity in the dataset. A machine cannot learn if there is no data available. Besides, a dataset with a lack of diversity gives the machine a hard time. A machine needs to have heterogeneity to learn meaningful insight. It is rare that an algorithm can extract information when there are no or few variations. It is recommended to have at least 20 observations per group to help the machine learn. This constraint leads to poor evaluation and prediction.

3.5.5 Application of Machine Learning

Augmentation

Machine learning, which assists humans with their day-to-day tasks, personally or commercially without having complete control of the output. Such machine learning is used in different ways such as Virtual Assistant, Data analysis, software solutions. The primary user is to reduce errors due to human bias.

Automation:

Machine learning, which works entirely autonomously in any field without the need for any human intervention. For example, robots performing the essential process steps in manufacturing plants.

Finance Industry

Machine learning is growing in popularity in the finance industry. Banks are mainly using ML to find patterns inside the data but also to prevent fraud.

Government organization

The government makes use of ML to manage public safety and utilities. Take the example of China with the massive face recognition. The government uses Artificial intelligence to prevent jaywalker.

Healthcare industry

Healthcare was one of the first industry to use machine learning with image detection.

Marketing

Broad use of AI is done in marketing thanks to abundant access to data. Before the age of mass data, researchers develop advanced mathematical tools like Bayesian analysis to estimate the value of a customer. With the boom of data, marketing department relies on AI to optimize the customer relationship and marketing campaign.

3.5.6 Example of application of Machine Learning in Supply Chain

Machine learning gives terrific results for visual pattern recognition, opening up many potential applications in physical inspection and maintenance across the entire supply chain network.

Unsupervised learning can quickly search for comparable patterns in the diverse dataset. In turn, the machine can perform quality inspection throughout the logistics hub, shipment with damage and wear.

For instance, IBM's Watson platform can determine shipping container damage. Watson combines visual and systems-based data to track, report and make recommendations in real-time.

In past year stock manager relies extensively on the primary method to evaluate and forecast the inventory. When combining big data and machine learning, better forecasting techniques have been implemented (an improvement of 20 to 30 % over traditional forecasting tools). In term of sales, it means an increase of 2 to 3 % due to the potential reduction in inventory costs.

3.5.7 Example of Machine Learning Google Car

For example, everybody knows the Google car. The car is full of lasers on the roof which are telling it where it is regarding the surrounding area. It has radar in the front, which is informing the car of the speed and motion of all the cars around it. It uses all of that data to figure out not only how to drive the car but also to figure out and predict what potential drivers around the car are going to do. What's impressive is that the car is processing almost a gigabyte a second of data.

Why is Machine Learning Important?

Machine learning is the best tool so far to analyze, understand and identify a pattern in the data. One of the main ideas behind machine learning is that the computer can be trained to automate tasks that would be exhaustive or impossible for a human being. The clear breach from the traditional analysis is that machine learning can take decisions with minimal human intervention.

Take the following example for this ML tutorial; a retail agent can estimate the price of a house based on his own experience and his knowledge of the market.

A machine can be trained to translate the knowledge of an expert into features. The features are all the characteristics of a house, neighborhood, economic environment, etc. that make the price difference. For the expert, it took him probably some years to master the art of estimate the price of a house. His expertise is getting better and better after each sale.

For the machine, it takes millions of data, (i.e., example) to master this art. At the very beginning of its learning, the machine makes a mistake, somehow like the junior salesman. Once the machine sees all the example, it got enough knowledge to make its estimation. At the same time, with incredible accuracy. The machine is also able to adjust its mistake accordingly.

Most of the big company have understood the value of machine learning and holding data. McKinsey have estimated that the value of analytics ranges from \$9.5 trillion to \$15.4 trillion while \$5 to 7 trillion can be attributed to the most advanced AI techniques.

Machine learning (ML) is the study of computer algorithms that improve automatically through experience. It is seen as a part of artificial intelligence. Machine learning algorithms build a model based on sample data, known as "training data", in order to make predictions or decisions without being explicitly programmed to do so. Machine learning algorithms are used in a wide variety of applications, such as email filtering and computer vision, where it is difficult or unfeasible to develop conventional algorithms to perform the needed tasks.

A subset of machine learning is closely related to computational statistics, which focuses on making predictions using computers; but not all machine learning is statistical learning. The study of mathematical optimization delivers methods, theory and application domains to the field of machine learning. Data mining is a related field of study, focusing on exploratory data analysis through unsupervised learning. In its application across business problems, machine learning is also referred to as predictive analytics.

Overview

Machine learning involves computers discovering how they can perform tasks without being explicitly programmed to do so. It involves computers learning from data provided so that they carry out certain tasks. For simple tasks assigned to computers, it is possible to program algorithms telling the machine how to execute all steps required to solve the problem at hand; on the computer's part, no learning is needed. For more advanced tasks, it can be challenging for a human to manually create the needed algorithms. In practice, it can turn out to be more effective to help the machine develop its own algorithm, rather than having human programmers specify every needed step.

Artificial intelligence

Machine Learning as subfield of AI. Part of Machine Learning as subfield of AI or part of AI as subfield of Machine Learning As a scientific endeavor, machine learning grew out of the quest for artificial intelligence. In the early days of AI as an academic discipline, some researchers were interested in having machines learn from data. They attempted to approach the problem with various symbolic methods, as well as what was then termed "neural networks"; these were mostly perceptrons and other models that were later found to be reinventions of the generalized linear models of statistics. Probabilistic reasoning was also employed, especially in automated medical diagnosis.

However, an increasing emphasis on the logical, knowledge-based approach caused a rift between AI and machine learning. Probabilistic systems were plagued by theoretical and practical problems of data acquisition and representation. By 1980, expert systems had come to dominate AI, and statistics was out of favor. Work on symbolic/knowledge-based learning did continue within AI, leading to inductive logic programming, but the more statistical line of research was now outside the field of AI proper, in pattern recognition and information retrieval. Neural networks research had been abandoned by AI and computer science around the same time. This line, too, was continued outside the AI/CS field, as "connectionism", by researchers from other disciplines including Hopfield, Rumelhart and Hinton. Their main success came in the mid-1980s with the reinvention of backpropagation.

Machine learning (ML), reorganized as a separate field, started to flourish in the 1990s. The field changed its goal from achieving artificial intelligence to tackling solvable problems of a practical nature. It shifted focus away from the symbolic approaches it had inherited from AI, and toward methods and models borrowed from statistics and probability theory.

As of 2020, many sources continue to assert that machine learning remains a subfield of AI. The main disagreement is whether all of ML is part of AI, as this would mean that anyone using ML could claim they are using AI. Others have the view that not all of ML is part of AI where only an 'intelligent' subset of ML is part of AI.

Data mining

Machine learning and data mining often employ the same methods and overlap significantly, but while machine learning focuses on prediction, based on known properties learned from the training data, data mining focuses on the discovery of (previously) unknown properties in the data (this is the analysis step of knowledge discovery in databases). Data mining uses many machine learning methods, but with different goals; on the other hand, machine learning also employs data mining methods as "unsupervised learning" or as a

preprocessing step to improve learner accuracy.

Much of the confusion between these two research communities (which do often have separate conferences and separate journals, ECML PKDD being a major exception) comes from the basic assumptions they work with: in machine learning, performance is usually evaluated with respect to the ability to reproduce known knowledge, while in knowledge discovery and data mining (KDD) the key task is the discovery of previously unknown knowledge. Evaluated with respect to known knowledge, an uninformed (unsupervised) method will easily be outperformed by other supervised methods, while in a typical KDD task, supervised methods cannot be used due to the unavailability of training data.

Optimization

Machine learning also has intimate ties to optimization: many learning problems are formulated as minimization of some loss function on a training set of examples. Loss functions express the discrepancy between the predictions of the model being trained and the actual problem instances (for example, in classification, one wants to assign a label to instances, and models are trained to correctly predict the pre-assigned labels of a set of examples).

Generalization

The difference between optimization and machine learning arises from the goal of generalization: while optimization algorithms can minimize the loss on a training set, machine learning is concerned with minimizing the loss on unseen samples. Characterizing the generalization of various learning algorithms is an active topic of current research, especially for deep learning algorithms.

Statistics

Machine learning and statistics are closely related fields in terms of methods, but distinct in their principal goal: statistics draws population inferences from a sample, while machine learning finds generalizable predictive patterns. According to Michael I. Jordan, the ideas of machine learning, from methodological principles to theoretical tools, have had a long pre-history in statistics. He also suggested the term data science as a placeholder to call the overall field.

Leo Breiman distinguished two statistical modeling paradigms: data model and algorithmic model, wherein "algorithmic model" means more or less the machine learning algorithms like Random forest.

Some statisticians have adopted methods from machine learning, leading to a combined field that they call statistical learning.

CHAPTER - 4

EXPERIMENTAL ANALYSIS AND DESIGN

Drowsy driving is a significant threat on the road, contributing to a substantial number of accidents. To address this issue, researchers are exploring various methods for driver drowsiness detection. This study investigates the application of Convolutional Neural Networks (CNN) for this critical task.

■ **Motivation:**

- Traditional drowsiness detection methods often rely on features like steering wheel movement or physiological signals, which may not be as robust under diverse driving conditions.
- CNN have demonstrated remarkable success in image recognition and classification tasks.
- By leveraging CNN, we aim to develop a system that can effectively detect driver drowsiness based on visual cues captured from in-vehicle cameras.

■ **Contribution of this Study:**

This research focuses on designing and evaluating a CNN-based system for driver drowsiness detection. We will explore:

■ **Network Architecture:**

Design and optimize a CNN architecture specifically tailored for driver drowsiness classification using facial features or other relevant visual cues.

■ **Data Acquisition and Preprocessing:**

Develop a strategy for collecting and pre-processing real-world or simulated driving data to train and validate the CNN model. This may involve techniques like data augmentation to handle variations in illumination and head pose.

■ **Performance Evaluation:**

Evaluate the performance of the proposed CNN model on unseen test data. Metrics such as accuracy, precision, recall, and F1-score will be used to assess the effectiveness of the system in detecting drowsy and alert driver states.

■ Expected Outcomes:

This research is expected to contribute to the development of a reliable and efficient driver drowsiness detection system using CNN. The findings can potentially improve road safety by providing real-time alerts to drivers who exhibit signs of drowsiness.

4.1 Data Flow Diagram

The DFD is also called as bubble chart. It is a simple graphical formalism that can be used to represent a system in terms of input data to the system, various processing carried out on this data, and the output data is generated by this system. The data flow diagram (DFD) is one of the most important modeling tools. It is used to model the system components. These components are the system process, the data used by the process, an external entity that interacts with the system and the information flows in the system.

DFD shows how the information moves through the system and how it is modified by a series of transformations. It is a graphical technique that depicts information flow and the transformations that are applied as data moves from input to output. DFD is also known as bubble chart. A DFD may be used to represent a system at any level of abstraction. DFD may be partitioned into levels that represent increasing information flow and functional detail. A Data Flow Diagram (DFD) is a graphical representation of how data flows through a system, depicting the processes that transform input data into output data. It consists of interconnected symbols that represent various components of the system and the flow of data between them.

Represented by rectangles, processes depict the functions or activities that manipulate input data to produce output data. Each process in the diagram describes a specific task or operation performed within the system. Represented by arrows, data flows depict the movement of data between processes, external entities, and data stores. They illustrate the path that data follows as it moves through the system, from its source to its destination. Represented by squares, external entities depict the sources or destinations of data that interact with the system but are external to it. These can include users, other systems, or external data sources. Represented by parallel lines, data stores depict the repositories where data is stored within the system. They represent persistent storage locations such as databases, files, or memory buffers.

DFD diagrams provide a high-level overview of the system's functionality and data flow, enabling stakeholders to understand the system's structure, identify potential bottlenecks or inefficiencies, and facilitate communication between stakeholders, designers, and developers. They are commonly used during the requirements analysis and design phases

of software development to visualize and document the flow of data within a system.

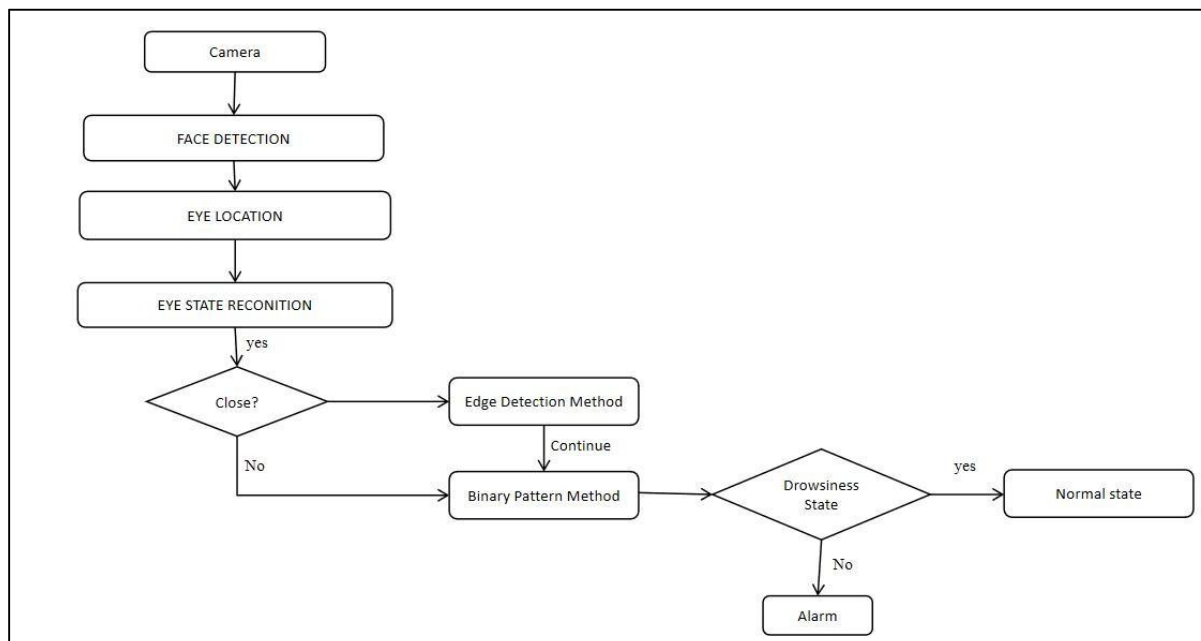


Figure 4.1 : DFD of STAY AWAKE STAY ALIVE

4.2 USE CASE DIAGRAM

A use case diagram is a visual representation of the interactions between users (actors) and a system, illustrating the various functionalities and behaviors of the system from the user's perspective. It consists of actors, use cases, and their relationships.

1. Actors: Actors are entities (human or system) that interact with the system to achieve specific goals. They are represented by stick figures and can be primary users, external systems, or other software components.

2. Use Cases: Use cases represent the functionalities or tasks that the system performs to fulfill the needs of its users. They are depicted as ovals and describe the system's behavior in response to user actions or events.

3. Relationships: Relationships between actors and use cases illustrate how actors interact with the system to accomplish specific tasks. Associations represent communication between actors and use cases, while include and extend relationships depict relationships between different use cases.

Use case diagrams provide a clear and concise overview of the system's functionality and its interactions with users, helping stakeholders understand the system's scope and requirements. They serve as a valuable tool for requirements analysis, design, and communication among project stakeholders.

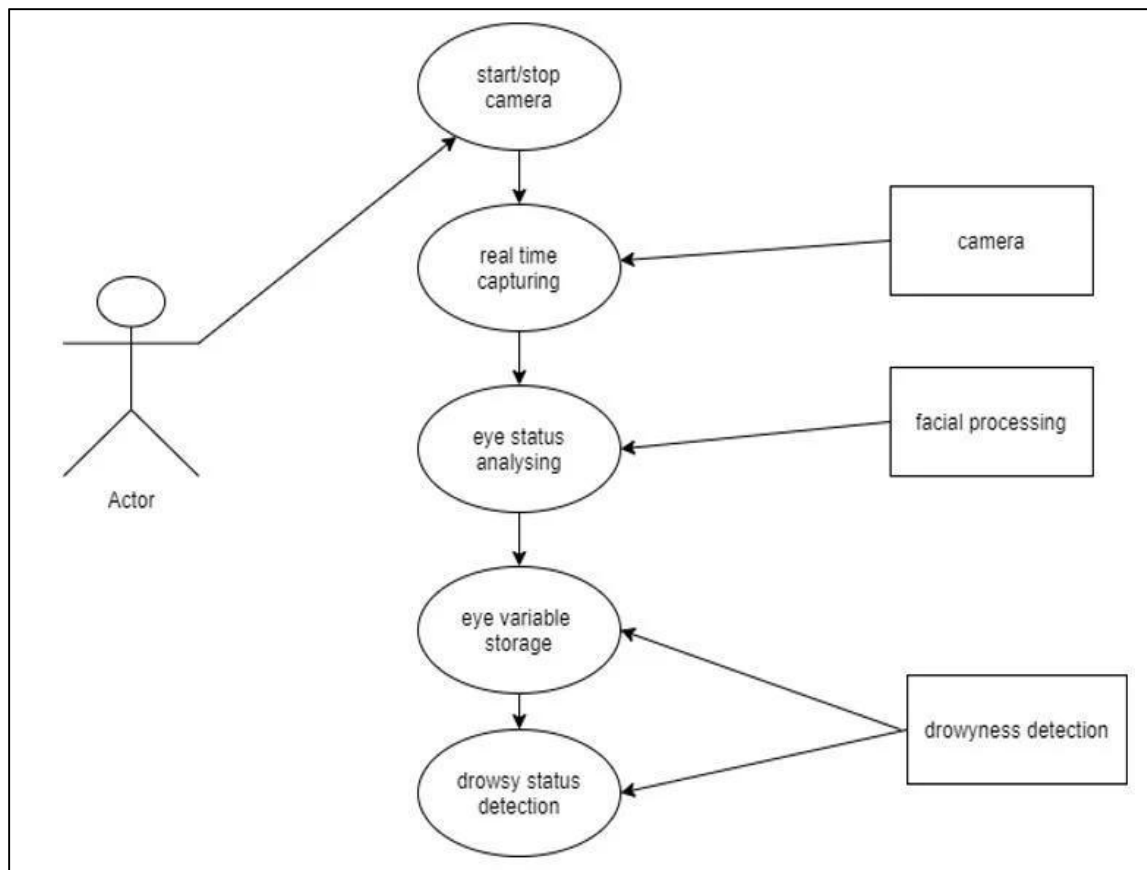


Figure 4.2 : Use-case Diagram

A use case diagram is a visual representation of how users interact with a system. It illustrates the various actions or tasks users perform within the system and the relationships between them. Actors, representing users or external systems, interact with the system through specific use cases, which depict individual functionality or features. Relationships such as associations, generalizations, and includes/extends dependencies are depicted to show how different elements interact. Use case diagrams aid in understanding system behavior, identifying system requirements, and communicating system functionality to stakeholders. They serve as a blueprint for system development, ensuring alignment between user needs and system design.

4.3 Sequence Diagram

A sequence diagram is a type of UML diagram that illustrates the interactions between objects or components in a system over time. It depicts the sequence of messages exchanged between objects to accomplish a specific task or scenario. Objects are represented as vertical lines, and messages are depicted as arrows between them, indicating the order and direction of communication. Sequence diagrams are valuable for visualizing the dynamic behavior of a system, identifying communication patterns, and understanding the flow of control during the execution of a use case or scenario.

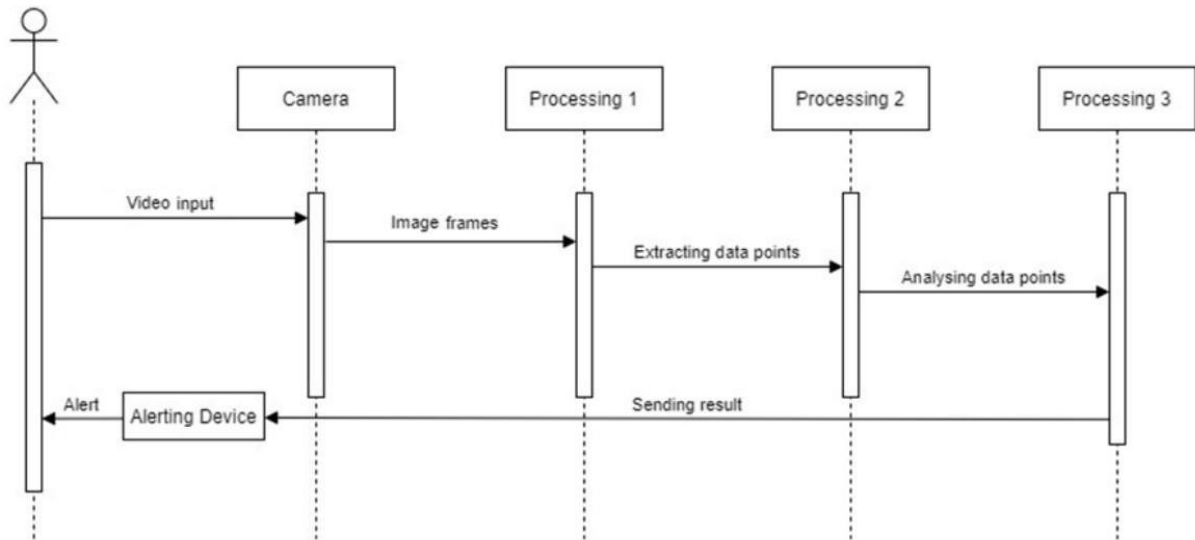


Figure 4.3 : Sequence Diagram

Long Short-Term Memory (LSTM) is a type of recurrent neural network (RNN) architecture that is specifically designed to address the vanishing gradient problem, which is a common issue in traditional RNNs. LSTMs are capable of learning long-term dependencies in sequential data by maintaining a memory cell that can store information over long periods of time.

The key components of an LSTM include:

1. **Cell State:** The cell state runs through the entire chain of LSTM units and acts as a conveyor belt, allowing information to flow unchanged. It enables LSTMs to retain information over long sequences, making them effective for processing sequential data.
2. **Forget Gate:** The forget gate is a sigmoid layer that decides what information from the cell state should be discarded or forgotten. It takes input from the previous LSTM unit and the current input and outputs a number between 0 and 1 for each element in the cell state, indicating how much of the information to keep.
3. **Input Gate:** The input gate is responsible for deciding what new information to store in the cell state. It consists of a sigmoid layer that determines which values will be updated and a tanh layer that creates a vector of new candidate values to be added to the cell state.
4. **Output Gate:** The output gate controls the information that is output from the cell state. It uses the current input and the previous output to decide what information should be passed on to the next LSTM unit or the output layer.

LSTMs have become popular in various applications, including natural language processing (NLP), speech recognition, time series forecasting, and more. Their ability to capture long-term dependencies in sequential data makes them well-suited for tasks that involve analyzing and generating sequences of data.

4.4 Results and Discussion

The existing systems, based on CNNs for real-time driver drowsiness detection, have shown considerable success in accurately identifying signs of drowsiness from facial expressions and eye movements. These systems utilize sophisticated CNN architectures to process facial images captured by in-car cameras, extracting features indicative of drowsiness such as eye closure, head nods, and changes in facial expressions. The results demonstrate that these CNN-based systems can effectively differentiate between alert and drowsy states, providing timely warnings to drivers to mitigate the risks of accidents caused by driver fatigue.

While the existing CNN-based drowsiness detection systems represent a significant advancement in automotive safety technology, there are certain limitations to consider. One such limitation is the reliance on facial images, which may be affected by factors such as lighting conditions, occlusions (e.g., glasses), and variations in facial expressions among different individuals. Additionally, although some systems incorporate advanced techniques like transfer learning and attention mechanisms to improve accuracy, there is still room for further enhancement, particularly in challenging scenarios where drivers exhibit subtle signs of drowsiness.

Aspects	Existing System	Proposed System
Focus	Primarily focuses on utilizing CNNs for analyzing facial expressions and eye movements for drowsiness detection.	Emphasizes integration of advanced driver assistance systems (ADAS) with CNNs for detecting drowsiness.
Methodology	Relies solely on CNNs for drowsiness detection.	Integrates visual cues from in-car monitoring systems (e.g., dashboard cameras, infrared sensors) with CNNs.
Approach	CNNs process facial images directly to detect drowsiness signs.	Utilizes visual cues extracted from in-car monitoring systems, enhancing the scope and accuracy of detection.
Data Source	Primarily relies on facial images captured by a camera pointed at the driver's face.	Utilizes visual data from in-car monitoring systems, which could include dashboard cameras or infrared sensors.
Potential Integration	Can be integrated into existing vehicle monitoring systems with facial recognition capabilities.	Can be integrated into existing ADAS frameworks, leveraging various in-car monitoring systems for data collection.

Enhancement	Subsequent systems might improve accuracy by incorporating additional features (e.g., eye gaze tracking).	Enhances detection accuracy by incorporating a broader range of visual cues from in-car monitoring systems.
Accuracy	97%	98.5%
Scope	Focused on real-time detection of drowsiness based on facial expressions and eye movements.	Expands the scope to include various visual cues from in-car monitoring systems for comprehensive detection.
Adaptability	Limited adaptability to diverse driving conditions and environments.	Potentially adaptable to various driving conditions and environments due to the integration of diverse data sources.
Potential Accuracy Improvement	May benefit from additional features or improved CNN architectures.	Expected to achieve higher accuracy through the utilization of diverse visual cues and improved system integration.

The proposed integration of CNNs with in-car monitoring systems aims to enhance drowsiness detection by leveraging a broader range of visual cues beyond facial image analysis alone. This integration offers several potential advantages, including the capture of additional contextual information such as vehicle speed, steering behavior, and environmental conditions.

Moreover, by incorporating data from multiple sources, including facial images, infrared images, and other sensor data, the proposed system enhances adaptability to diverse driving conditions and environments. However, the implementation of the proposed system may pose challenges related to data integration, system complexity, and real-time processing requirements. Ensuring seamless integration with existing ADAS frameworks and optimizing computational efficiency will be crucial for practical deployment in commercial vehicles.

In conclusion, both the existing CNN-based drowsiness detection systems and the proposed integration of CNNs with in-car monitoring systems hold promise for advancing the effectiveness of driver drowsiness detection technology. By leveraging a combination of visual cues and contextual information, these approaches have the potential to enhance road safety and prevent accidents caused by impaired driving. Further research and development efforts are needed to validate the efficacy of these approaches in real-world driving scenarios and address practical implementation challenges.




Figure 4.4 : awake result




Figure 4.5 : Drowsiness result

4.5 Requirement Analysis

Requirement analysis, also called requirement engineering, is the process of determining user expectations for a new modified product. It encompasses the tasks that determine the need for analyzing, documenting, validating and managing software or system requirements. The requirements should be documentable, actionable, measurable, testable and traceable related to identified business needs or opportunities and define to a level of detail, sufficient for system design.

FUNCTIONAL REQUIREMENTS

It is a technical specification requirement for the software products. It is the first step in the requirement analysis process which lists the requirements of particular software systems including functional, performance and security requirements. The function of the system depends mainly on the quality hardware used to run the software with given functionality.

Usability

It specifies how easy the system must be use. It is easy to ask queries in any format which is short or long, porter stemming algorithm stimulates the desired response for user.

Robustness

It refers to a program that performs well not only under ordinary conditions but also under unusual conditions. It is the ability of the user to cope with errors for irrelevant queries during execution.

Security

The state of providing protected access to resource is security. The system provides good security and unauthorized users cannot access the system there by providing high security.

Reliability

It is the probability of how often the software fails. The measurement is often expressed in MTBF (Mean Time Between Failures). The requirement is needed in order to ensure that the processes work correctly and completely without being aborted. It can handle any load and survive and even capable of working around any failure.

Compatibility

It is supported by version above all web browsers. Using any web servers like local host makes the system real-time experience.

Flexibility

The flexibility of the project is provided in such a way that it has the ability to run on different environments being executed by different users.

Safety

Safety is a measure taken to prevent trouble. Every query is processed in a secured manner without letting others to know one's personal information.

NON- FUNCTIONAL REQUIREMENTS

Portability

It is the usability of the same software in different environments. The project can be run in any operating system.

Performance

These requirements determine the resources required, time interval, throughput and everything that deals with the performance of the system.

Accuracy

The result of the requesting query is very accurate and high speed of retrieving information. The degree of security provided by the system is high and effective.

Maintainability

Project is simple as further updates can be easily done without affecting its stability. Maintainability basically defines that how easy it is to maintain the system. It means that how easy it is to maintain the system, analyse, change and test the application. Maintainability of this project is simple as further updates can be easily done without affecting its stability.

4.6 System Design And Testing Plan

4.6.1 IMAGE PROCESSING

What is Image Processing?

Image processing is a method to convert an image into digital form and perform some operations on it, in order to get an enhanced image or to extract some useful information from it. It is a type of signal dispensation in which input is image, like video frame or photograph and output may be image or characteristics associated with that image. Usually Image Processing system includes treating images as two dimensional signals while applying already set signal processing methods to them.

It is among rapidly growing technologies today, with its applications in various aspects of a business. Image Processing forms core research area within engineering and computer science disciplines too.

Image processing basically includes the following three steps.

- Importing the image with optical scanner or by digital photography.
- Analyzing and manipulating the image which includes data compression and image enhancement and spotting patterns that are not to human eyes like satellite photographs.
- Output is the last stage in which result can be altered image or report that is based on image analysis.

Purpose of Image processing

The purpose of image processing is divided into 5 groups. They are:

1. Visualization - Observe the objects that are not visible.
2. Image sharpening and restoration - To create a better image.
3. Image retrieval - Seek for the image of interest.
4. Measurement of pattern – Measures various objects in an image.
5. Image Recognition – Distinguish the objects in an image.

Types

The two types of **methods used for Image Processing** are **Analog and Digital** Image Processing. Analog or visual techniques of image processing can be used for the hard copies like printouts and photographs. Image analysts use various fundamentals of interpretation while using these visual techniques. The image processing is not just confined to area that has to be studied but on knowledge of analyst. Association is another important tool in image processing through visual techniques. So analysts apply a combination of personal knowledge and collateral data to image processing.

Digital Processing techniques help in manipulation of the digital images by using

computers. As raw data from imaging sensors from satellite platform contains deficiencies. To get over such flaws and to get originality of information, it has to undergo various phases of processing. The three general phases that all types of data have to undergo while using digital technique are Pre- processing, enhancement and display, information extraction.

IMAGE PROCESSING CONCEPTS

Binary Images

Binary images are images whose pixels have only two possible intensity values. They are normally displayed as black and white. Numerically, the two values are often 0 for black, and either 1 or 255 for white.

Binary images are often produced by thresholding a grayscale or color image, in order to separate an object in the image from the background. The color of the object (usually white) is referred to as the foreground color. The rest (usually black) is referred to as the background color. However, depending on the image which is to be thresholded, this polarity might be inverted, in which case the object is displayed with 0 and the background is with a non-zero value.

Some morphological operators assume a certain polarity of the binary input image so that if we process an image with inverse polarity the operator will have the opposite effect. For example, if we apply a closing operator to a black text on white background, the text will be opened.

Color Images

It is possible to construct (almost) all visible colors by combining the three primary colors red, green and blue, because the human eye has only three different color receptors, each of them sensible to one of the three colors. Different combinations in the stimulation of the receptors enable the human eye to distinguish approximately 350000 colors. A RGB color image is a multi-spectral image with one band for each color red, green and blue, thus producing a weighted combination of the three primary colors for each pixel.

A full 24-bit color image contains one 8-bit value for each color, thus being able to display $2^{24}=16777216$ different colors.

However, it is computationally expensive and often not necessary to use the full 24-bit image to store the color for each pixel. Therefore, the color for each pixel is often encoded in a single byte, resulting in an 8-bit color image. The process of reducing the color

representation from 24-bits to 8-bits, known as color quantization, restricts the number of possible colors to 256. However, there is normally no visible difference between a 24-color image and the same image displayed with 8 bits. An 8-bit color images are based on colormaps, which are look-up tables taking the 8-bit pixel value as index and providing an output value for each color.

8-bit Color Images

Full RGB color requires that the intensities of three color components be specified for each and every pixel. It is common for each component intensity to be stored as an 8-bit integer, and so each pixel requires 24 bits to completely and accurately specify its color. If this is done, then the image is known as a 24-bit color image. However there are two problems with this approach:

- Storing 24 bits for every pixel leads to very large image files that with current technology are cumbersome to store and manipulate. For instance a 24-bit 512×512 image takes up 750KB in uncompressed form.
- Many monitor displays use colormaps with 8-bit index numbers, meaning that they can only display 256 different colors at any one time. Thus it is often wasteful to store more than 256 different colors in an image anyway, since it will not be possible to display them all on screen.

Because of this, many image formats (e.g. 8-bit GIF and TIFF) use 8-bit colormaps to restrict the maximum number of different colors to 256. Using this method, it is only necessary to store an 8-bit index into the colormap for each pixel, rather than the full 24-bit color value. Thus 8-bit image formats consist of two parts: a colormap describing what colors are present in the image, and the array of index values for each pixel in the image.

When a 24-bit full color image is turned into an 8-bit image, it is usually necessary to throw away some of the colors, a process known as color quantization. This leads to some degradation in image quality, but in practice the observable effect can be quite small, and in any case, such degradation is inevitable if the image output device (e.g. screen or printer) is only capable of displaying 256 colors or less.

The use of 8-bit images with colormaps does lead to some problems in image processing. First of all, each image has to have its own colormap, and there is usually no guarantee that each image will have exactly the same colormap. Thus on 8-bit displays it is frequently impossible to correctly display two different color images that have different colormaps at the same time. Note that in practice 8-bit images often use reduced size

colormaps with less than 256 colors in order to avoid this problem.

Another problem occurs when the output image from an image processing operation contains different colors to the input image or images. This can occur very easily, as for instance when two color images are added together pixel-by-pixel. Since the output image contains different colors from the input images, it ideally needs a new colormap, different from those of the input images, and this involves further color quantization which will degrade the image quality. Hence the resulting output is usually only an approximation of the desired output. Repeated image processing operations will continually degrade the image colors. And of course we still have the problem that it is not possible to display the images simultaneously with each other on the same 8-bit display.

Because of these problems it is to be expected that as computer storage and processing power become cheaper, there will be a shift away from 8-bit images and towards full 24-bit image processing.

24-bit Color Images

Full RGB color requires that the intensities of three color components be specified for each and every pixel. It is common for each component intensity to be stored as an 8-bit integer, and so each pixel requires 24 bits to completely and accurately specify its color. Image formats that store a full 24 bits to describe the color of each and every pixel are therefore known as 24-bit color images.

Using 24 bits to encode color information allows $2^{24}=16777216$ different colors to be represented, and this is sufficient to cover the full range of human color perception fairly well.

The term 24-bit is also used to describe monitor displays that use 24 bits per pixel in their display memories, and which are hence capable of displaying a full range of colors.

There are also some disadvantages to using 24-bit images. Perhaps the main one is that it requires three times as much memory, disk space and processing time to store and manipulate 24-bit color images as compared to 8-bit color images. In addition, there is often not much point in being able to store all those different colors if the final output device (e.g. screen or printer) can only actually produce a fraction of them. Since it is possible to use colormaps to produce 8-bit color images that look almost as good, at the time of writing 24-bit displays are relatively little used. However it is to be expected that as the technology becomes

Color Quantization

Color quantization is applied when the color information of an image is to be reduced. The most common case is when a 24-bit color image is transformed into an 8-bit color image.

Two decisions have to be made:

1. which colors of the larger color set remain in the new image, and
2. how are the discarded colors mapped to the remaining ones.

The simplest way to transform a 24-bit color image into 8 bits is to assign 3 bits to red and green and 2 bits to blue (blue has only 2 bits, because of the eye's lower sensitivity to this color). This enables us to display 8 different shades of red and green and 4 of blue. However, this method can yield only poor results. For example, an image might contain different shades of blue which are all clustered around a certain value such that only one shade of blue is used in the 8-bit image and the remaining three blues are not used.

Alternatively, since 8-bit color images are displayed using a colormap, we can assign any arbitrary color to each of the 256 8-bit values and we can define a separate colormap for each image. This enables us perform a color quantization adjusted to the data contained in the image. One common approach is the popularity algorithm, which creates a histogram of all colors and retains the 256 most frequent ones. Another approach, known as the median-cut algorithm, yields even better results but also needs more computation time. This technique recursively fits a box around all colors used in the RGB colorspace which it splits at the median value of its longest side. The algorithm stops after 255 recursions. All colors in one box are mapped to the centroid of this box.

All above techniques restrict the number of displayed colors to 256. A technique of achieving additional colors is to apply a variation of half-toning used for gray scale images, thus increasing the color resolution at the cost of spatial resolution. The 256 values of the colormap are divided into four sections containing 64 different values of red, green, blue and white. As can be seen in Figure 1, a 2×2 pixel area is grouped together to represent one composite color, each of the four pixels displays either one of the primary colors or white. In this way, the number of possible colors is increased from 256 to 64^4 .




Figure 4.6 : A 2×2 pixel area displaying one composite color

Convolution

Convolution is a simple mathematical operation which is fundamental to many common image processing operators. Convolution provides a way of 'multiplying together' two arrays of numbers, generally of different sizes, but of the same dimensionality, to produce a third array of numbers of the same dimensionality. This can be used in image processing to implement operators whose output pixel values are simple linear combinations of certain input pixel values.

Convolutional Neural Network (CNN) Architecture for driver drowsy detection

Introduction

The CNN architecture designed for driver drowsy detection utilizes a series of convolutional and dense layers to effectively extract features and classify digits from input images. This architecture is specifically tailored to handle the complexities and variations present in handwritten digits.

Convolutional Layers

The CNN begins with three convolutional layers, each consisting of multiple filters that convolve across the input images. These filters extract hierarchical features such as edges, corners, and textures from the input digits. The depth of the filters increases with the depth of the network, allowing the model to capture increasingly abstract representations of the input data.

Pooling Layers

Following each convolutional layer, max-pooling layers are employed to down

sample the feature maps, reducing their spatial dimensions while retaining important information. This process helps in creating translation-invariant representations and reduces the computational complexity of subsequent layers.

Dense Layers

After the convolutional and pooling layers, the feature maps are flattened and passed through three dense layers. These dense layers serve as classifiers, transforming the extracted features into predictions for each digit class. The number of neurons in these layers gradually decreases, leading to a compact representation of the digit features.

Activation Functions

Throughout the network, Rectified Linear Unit (ReLU) activation functions are applied after each convolutional and dense layer. ReLU introduces non-linearity into the model, enabling it to learn complex patterns and relationships within the data.

Optimizer and Validation

The Adam optimizer is utilized to optimize the network's weights during training. Adam combines the advantages of both AdaGrad and RMSProp, providing efficient optimization and convergence. Additionally, to evaluate the model's performance and prevent overfitting, a validation set consisting of 10,000 data points is used during training.

Training Procedure

During training, the model iteratively adjusts its parameters to minimize the difference between predicted and actual labels using backpropagation. The training data is fed through the network in batches, and the Adam optimizer updates the weights based on the computed gradients.

Evaluation and Performance

After training, the model's performance is evaluated using various metrics such as accuracy, precision, recall, and F1-score. The validation set is used to assess the model's generalization ability and identify any potential issues such as overfitting or underfitting.

In conclusion, the CNN architecture for handwritten digit recognition leverages the power of convolutional and dense layers to effectively learn and classify digit representations from input images. Through proper optimization, validation, and training procedures, the model can achieve high accuracy and robustness in recognizing handwritten digits across diverse styles and variations.




Fig 4.7 - Convolutional Neural Network

In an image processing context, one of the input arrays is normally just a grayscale image. The second array is usually much smaller, and is also two-dimensional (although it may be just a single pixel thick), and is known as the kernel. Figure 1 shows an example image and kernel that we will use to illustrate convolution.

I₁₁	I₁₂	I₁₃	I₁₄	I₁₅	I₁₆	I₁₇	I₁₈	I₁₉	<table><tr><td>K₁₁</td><td>K₁₂</td><td>K₁₃</td></tr><tr><td>K₂₁</td><td>K₂₂</td><td>K₂₃</td></tr></table>	K₁₁	K₁₂	K₁₃	K₂₁	K₂₂	K₂₃
K₁₁	K₁₂	K₁₃													
K₂₁	K₂₂	K₂₃													
I₂₁	I₂₂	I₂₃	I₂₄	I₂₅	I₂₆	I₂₇	I₂₈	I₂₉							
I₃₁	I₃₂	I₃₃	I₃₄	I₃₅	I₃₆	I₃₇	I₃₈	I₃₉							
I₄₁	I₄₂	I₄₃	I₄₄	I₄₅	I₄₆	I₄₇	I₄₈	I₄₉							
I₅₁	I₅₂	I₅₃	I₅₄	I₅₅	I₅₆	I₅₇	I₅₈	I₅₉							
I₆₁	I₆₂	I₆₃	I₆₄	I₆₅	I₆₆	I₆₇	I₆₈	I₆₉							

Figure 4.8 : An example small image (left) and kernel (right) to illustrate convolution.

The convolution is performed by sliding the kernel over the image, generally starting at the top left corner, so as to move the kernel through all the positions where the kernel fits entirely within the boundaries of the image. (Note that implementations differ in what they do at the edges of images, as explained below.) Each kernel position corresponds to a single output pixel, the value of which is calculated by multiplying together the kernel value and the underlying image pixel value for each of the cells in the kernel, and then adding all these numbers together.

So, in our example, the value of the bottom right pixel in the output image will be given by:

$$O_{57} = I_{57}K_{11} + I_{58}K_{13} + I_{59}K_{13} + I_{67}K_{21} + I_{68}K_{22} + I_{69}K_{23}$$

If the image has M rows and N columns, and the kernel has m rows and n columns, then the size of the output image will have $M - m + 1$ rows, and $N - n + 1$ columns.

Mathematically we can write the convolution as:

$$O(i, j) = \sum_{k=1}^m \sum_{l=1}^n I(i + k - 1, j + l - 1)K(k, l)$$

where i runs from 1 to $M - m + 1$ and j runs from 1 to $N - n + 1$.

Note that many implementations of convolution produce a larger output image than this because they relax the constraint that the kernel can only be moved to positions where it fits entirely within the image. Instead, these implementations typically slide the kernel to all positions where just the top left corner of the kernel is within the image. Therefore the kernel 'overlaps' the image on the bottom and right edges. One advantage of this approach is that the output image is the same size as the input image. Unfortunately, in order to calculate the output pixel values for the bottom and right edges of the image, it is necessary to invent input pixel values for places where the kernel extends off the end of the image. Typically pixel values of zero are chosen for regions outside the true image, but this can often distort the output image at these places. Therefore in general if you are using a convolution implementation that does this, it is better to clip the image to remove these spurious regions. Removing $n - 1$ pixels from the right hand side and $m - 1$ pixels from the bottom will fix things. Convolution can be used to implement many different operators, particularly spatial filters and feature detectors. Examples include Gaussian smoothing and the Sobel edge detector.

Distance Metrics

It is often useful in image processing to be able to calculate the distance between two pixels in an image, but this is not as straightforward as it seems. The presence of the pixel grid makes several so-called distance metrics possible which often give different answers to each other for the distance between the same pair of points. We consider the three most important ones.

Euclidean Distance

This is the familiar straight line distance that most people are familiar with. If the two pixels that we are considering have coordinates (x_1, y_1) , (x_2, y_2) then the Euclidean distance is given by:

$$D_{\text{Euclid}} = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$

City Block Distance

Also known as the Manhattan distance. This metric assumes that in going from one pixel to the other it is only possible to travel directly along pixel grid lines. Diagonal moves are not allowed. Therefore the 'city block' distance is given by:

$$D_{\text{City}} = |x_2 - x_1| + |y_2 - y_1|$$

Chessboard Distance

This metric assumes that you can make moves on the pixel grid as if you were a King making moves in chess, *i.e.* a diagonal move counts the same as a horizontal move. This means that the metric is given by:

$$D_{\text{Chess}} = \max(|x_2 - x_1|, |y_2 - y_1|)$$

Note that the last two metrics are usually much faster to compute than the Euclidean metric and so are sometimes used where speed is critical but accuracy is not too important.

Dithering

Dithering is an image display technique that is useful for overcoming limited display resources. The word dither refers to a random or semi-random perturbation of the pixel values.

Two applications of this techniques are particularly useful:

Low quantization display: When images are quantized to a few bits (e.g. 3) then only a limited number of graylevels are used in the display of the image. If the scene is smoothly shaded, then the image display will generate rather distinct boundaries around the edges of image regions when the original scene intensity moves from one quantization level to the next. To eliminate this effect, one dithering technique adds random noise (with a small range of values) to the input signal before quantization into the output range. This randomizes the quantization of the pixels at the original quantization boundary, and thus pixels make a more

gradual transition from neighborhoods containing 100% of the first quantization level to neighborhoods containing 100% of the second quantization level.

Limited color display: When fewer colors are able to be displayed (e.g. 256) than are present in the input image (e.g. 24 bit color), then patterns of adjacent pixels are used to simulate the appearance of the unrepresented colors.

Edge Detectors

Edges are places in the image with strong intensity contrast. Since edges often occur at image locations representing object boundaries, edge detection is extensively used in image segmentation when we want to divide the image into areas corresponding to different objects. Representing an image by its edges has the further advantage that the amount of data is reduced significantly while retaining most of the image information.

Since edges consist of mainly high frequencies, we can, in theory, detect edges by applying a highpass frequency filter in the Fourier domain or by convolving the image with an appropriate kernel in the spatial domain. In practice, edge detection is performed in the spatial domain, because it is computationally less expensive and often yields better results. Representing an image by its edges has the further advantage that the amount of data is reduced significantly while retaining most of the image information.

Since edges consist of mainly high frequencies, we can, in theory, detect edges by applying a highpass frequency filter in the Fourier domain or by convolving the image with an appropriate kernel in the spatial domain

Since edges correspond to strong illumination gradients, we can highlight them by calculating the derivatives of the image. This is illustrated for the one-dimensional case in

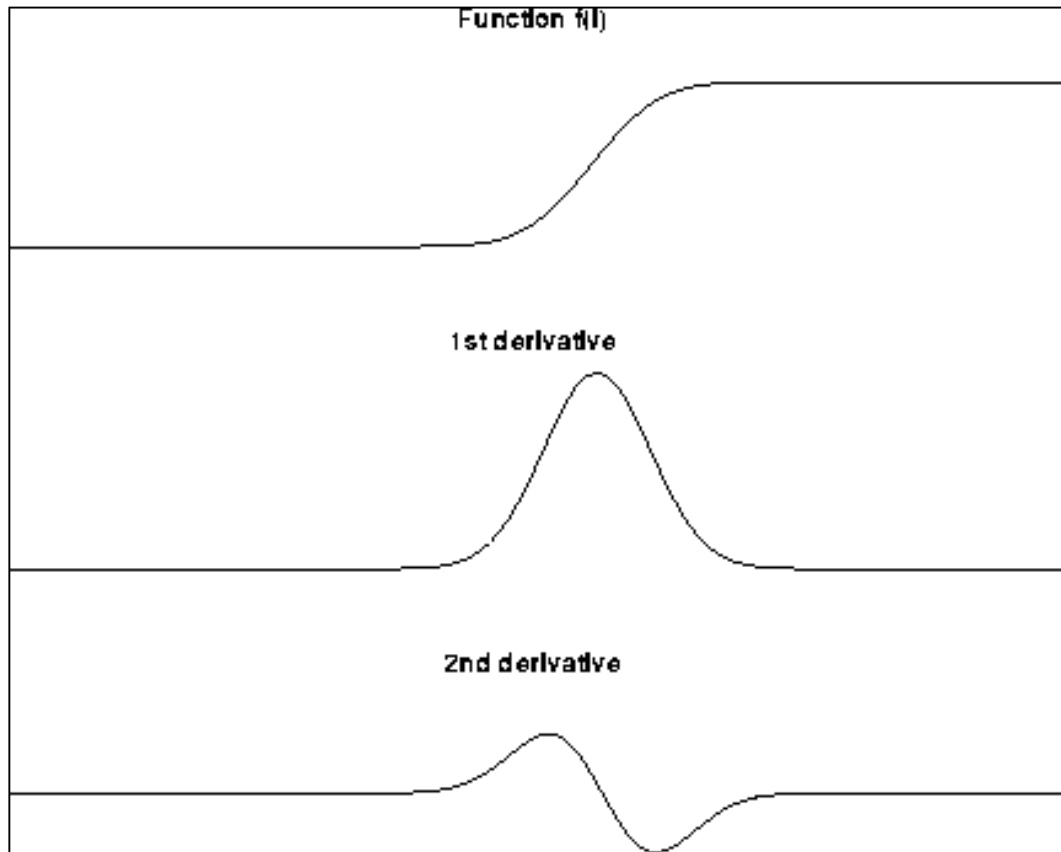


Figure 4.9 1st and 2nd derivative of an edge illustrated in one dimension.

We can see that the position of the edge can be estimated with the maximum of the 1st derivative or with the zero-crossing of the 2nd derivative. Therefore we want to find a technique to calculate the derivative of a two-dimensional image. For a discrete one-dimensional function $f(i)$, the first derivative can be approximated by

$$\frac{df(i)}{d(i)} = f(i+1) - f(i)$$

Calculating this formula is equivalent to convolving the function with $[-1 \ 1]$. Similarly the 2nd derivative can be estimated by convolving $f(i)$ with $[1 \ -2 \ 1]$.

Different edge detection kernels which are based on the above formula enable us to calculate either the 1st or the 2nd derivative of a two-dimensional image. There are two common approaches to estimate the 1st derivative in a two-dimensional image, Prewitt compass edge detection and gradient edge detection.

Prewitt compass edge detection involves convolving the image with a set of (usually 8) kernels, each of which is sensitive to a different edge orientation. The kernel producing the maximum response at a pixel location determines the edge magnitude and orientation. Different sets of kernels might be used: examples include Prewitt, Sobel, Kirsch and Robinson kernels.

Gradient edge detection is the second and more widely used technique. Here, the image is convolved with only two kernels, one estimating the gradient in the x-direction, G_x , the other the gradient in the y-direction, G_y . The absolute gradient magnitude is then given by

$$|G| = \sqrt{G_x^2 + G_y^2}$$

and is often approximated with

$$|G| = |G_x| + |G_y|$$

In many implementations, the gradient magnitude is the only output of a gradient edge detector, however the edge orientation might be calculated with

The most common kernels used for the gradient edge detector are the Sobel, Roberts Cross and Prewitt operators.

After having calculated the magnitude of the 1st derivative, we now have to identify those pixels corresponding to an edge. The easiest way is to threshold the gradient image, assuming that all pixels having a local gradient above the threshold must represent an edge. An alternative technique is to look for local maxima in the gradient image, thus producing one pixel wide edges. A more sophisticated technique is used by the Canny edge detector. It first applies a gradient edge detector to the image and then finds the edge pixels using non-maximal suppression and hysteresis tracking.

An operator based on the 2nd derivative of an image is the Marr edge detector, also known as zero crossing detector. Here, the 2nd derivative is calculated using a Laplacian of Gaussian (LoG) filter. The Laplacian has the advantage that it is an isotropic measure of the 2nd derivative of an image, i.e. the edge magnitude is obtained independently from the edge orientation by convolving the image with only one kernel. The edge positions are then given by the zero-crossings in the LoG image. The scale of the edges which are to be detected can be controlled by changing the variance of the Gaussian.

A general problem for edge detection is its sensitivity to noise, the reason being that calculating the derivative in the spatial domain corresponds to accentuating high frequencies and hence magnifying noise. This problem is addressed in the Canny and Marr operators by convolving the image with a smoothing operator (Gaussian) before calculating the derivative.

Frequency Domain

For simplicity, assume that the image I being considered is formed by projection from scene S (which might be a two- or three-dimensional scene, etc.).

The frequency domain is a space in which each image value at image position F represents the amount that the intensity values in image I vary over a specific distance related to F . In the frequency domain, changes in image position correspond to changes in the spatial frequency, (or the rate at which image intensity values) are changing in the spatial domain image I .

For example, suppose that there is the value 20 at the point that represents the frequency 0.1 (or 1 period every 10 pixels). This means that in the corresponding spatial domain image I the intensity values vary from dark to light and back to dark over a distance of 10 pixels, and that the contrast between the lightest and darkest is 40 gray levels (2 times 20).

The spatial frequency domain is interesting because: 1) it may make explicit periodic relationships in the spatial domain, and 2) some image processing operators are more efficient or indeed only practical when applied in the frequency domain.

In most cases, the Fourier Transform is used to convert images from the spatial domain into the frequency domain and vice-versa.

A related term used in this context is spatial frequency, which refers to the (inverse of the) periodicity with which the image intensity values change. Image features with high spatial frequency (such as edges) are those that change greatly in intensity over short image distances.

Grayscale Images

A grayscale (or graylevel) image is simply one in which the only colors are shades of gray. The reason for differentiating such images from any other sort of color image is that less information needs to be provided for each pixel. In fact a 'gray' color is one in which the red, green and blue components all have equal intensity in RGB space, and so it is only necessary to specify a single intensity value for each pixel, as opposed to the three intensities needed to specify each pixel in a full color image.

Often, the grayscale intensity is stored as an 8-bit integer giving 256 possible different

shades of gray from black to white. If the levels are evenly spaced then the difference between successive graylevels is significantly better than the graylevel resolving power of the human eye.

Grayscale images are very common, in part because much of today's display and image capture hardware can only support 8-bit images. In addition, grayscale images are entirely sufficient for many tasks and so there is no need to use more complicated and harder-to-process color images.

Image Editing Software

There is a huge variety of software for manipulating images in various ways. Much of this software can be grouped under the heading image processing software, and the bulk of this reference is concerned with that group.

Another very important category is what we call image editing software. This group includes painting programs, graphic art packages and so on. They are often useful in conjunction with image processing software packages, in situations where direct immediate interaction with an image is the easiest way of achieving something. For instance, if a region of an image is to be masked out for subsequent image processing, it may be easiest to create the mask using an art package by directly drawing on top of the original image. The mask used in the description of the AND operator was created this way for instance. Art packages also often allow the user to move sections of the images around and brighten or darken selected regions interactively. Few dedicated image processing packages offer the same flexibility and ease of use in this respect.

Isotropic Operators

An isotropic operator in an image processing context is one which applies equally well in all directions in an image, with no particular sensitivity or bias towards one particular set of directions (e.g. compass directions). A typical example is the zero crossing edge detector which responds equally well to edges in any orientation. Another example is Gaussian smoothing. It should be borne in mind that although an operator might be isotropic in theory, the actual implementation of it for use on a discrete pixel grid may not be perfectly isotropic. An example of this is a Gaussian smoothing filter with very small standard deviation on a square grid.

Kernel

A kernel is a (usually) smallish matrix of numbers that is used in image convolutions. Differently sized kernels containing different patterns of numbers give rise to different results under convolution. For instance, Figure 1 shows a 3×3 kernel that implements a mean filter.

1	1	1
1	1	1
1	1	1

Set of coordinate points =
{ (-1, -1), (0, -1), (1, -1),
(-1, 0), (0, 0), (1, 0),
(-1, 1), (0, 1), (1, 1) }

Figure 4.10 : Convolution kernel for a mean filter with 3×3 neighborhood.

The word 'kernel' is also commonly used as a synonym for 'structuring element', which is a similar object used in mathematical morphology. A structuring element differs from a kernel in that it also has a specified origin. This sense of the word 'kernel' is not used in HIPR.

Logical Operators

Logical operators are generally derived from Boolean algebra, which is a mathematical way of manipulating the truth values of concepts in an abstract way without bothering about what the concepts actually mean. The truth value of a concept in Boolean value can have just one of two possible values: true or false. Boolean algebra allows you to represent things like: $A \text{ AND } B$

where A represents 'The block is red', and B represents 'The block is large'. Now each of these sub-phrases has its own truth value in any given situation: each sub-phrase is either true or false. Moreover, the entire composite phrase also has a truth value: it is true if both of the sub-phrases are true, and false in any other case. We can write this AND combination rule (and its dual operation NAND) using a truth-table as shown in Figure 1, in which we conventionally represent true by 1, and false by zero.

A	B	Q	A	B	Q
0	0	0	0	0	1
0	1	0	0	1	1
1	0	0	1	0	1
1	1	1	1	1	0
AND			NAND		

Figure 4.11: Truth-tables for AND and NAND

The left hand table shows each of the possible combinations of truth values of A and B, and the the resulting truth value of A AND B. Similar truth-tables can be set up for the other logical operators: NAND, OR, NOR, XOR, XNOR and NOT.

Turning now to an image processing context, the pixel values in a binary image, which are either 0 or 1, can be interpreted as truth values as above. Using this convention we can carry out logical operations on images simply by applying the truth-table combination rules to the pixel values from a pair of input images (or a single input image in the case of NOT). Normally, corresponding pixels from each of two identically sized binary input images are compared to produce the output image, which is another binary image of the same size. As with other image arithmetic operations, it is also possible to logically combine a single input image with a constant logical value, in which case each pixel in the input image is compared to the same constant in order to produce the corresponding output pixel. See the individual logical operator descriptions for examples of these operations.

Logical operations can also be carried out on images with integer pixel values. In this extension the logical operations are normally carried out in bitwise fashion on binary representations of those integers, comparing corresponding bits with corresponding bits to produce the output pixel value. For instance, suppose that we wish to XOR the integers 47 and 255 together using 8-bit integers. 47 is 00101111 in binary and 255 is 11111111.

XORing these together in bitwise fashion, we have 11010000 in binary or 208 in decimal.

Note that not all implementations of logical operators work in such bitwise fashion. For instance some will treat zero as false and any non-zero value as true and will then apply the conventional 1-bit logical functions to derive the output image. The output may be a simple binary image itself, or it may be a graylevel image formed perhaps by multiplying what would be the binary output image (containing 0's and 1's) with one of the input images.

Look-up Tables and Colormaps

Look-Up Tables or LUTs are fundamental to many aspects of image processing. An LUT is simply a table of cross-references linking index numbers to output values. The most common use is to determine the colors and intensity values with which a particular image will be displayed, and in this context the LUT is often called simply a colormap.

The idea behind the colormap is that instead of storing a definite color for each pixel in an image, for instance in 24-bit RGB format, each pixel's value is instead treated as an index number into the colormap. When the image is to be displayed or otherwise processed, the colormap is used to look up the actual colors corresponding to each index number. Typically, the output values stored in the LUT would be RGB color values.

There are two main advantages to doing things this way. Firstly, the index number can be made to use fewer bits than the output value in order to save storage space. For instance an 8-bit index number can be used to look up a 24-bit RGB color value in the LUT. Since only the 8-bit index number needs to be stored for each pixel, such 8-bit color images take up less space than a full 24-bit image of the same size. Of course the image can only contain 256 different colors (the number of entries in an 8-bit LUT), but this is sufficient for many applications and usually the observable image degradation is small.

Secondly the use of a color table allows the user to experiment easily with different color labeling schemes for an image.

One disadvantage of using a colormap is that it introduces additional complexity into an image format. It is usually necessary for each image to carry around its own colormap, and this LUT must be continually consulted whenever the image is displayed or processed.

Another problem is that in order to convert from a full color image to (say) an 8-bit color image using a color image, it is usually necessary to throw away many of the original colors, a process known as color quantization. This process is lossy, and hence the image

quality is degraded during the quantization process. Additionally, when performing further image processing on such images, it is frequently necessary to generate a new colormap for the new images, which involves further color quantization, and hence further image degradation.

As well as their use in colormaps, LUTs are often used to remap the pixel values within an image. This is the basis of many common image processing point operations such as thresholding, gamma correction and contrast stretching. The process is often referred to as anamorphosis.

Masking

A mask is a binary image consisting of zero- and non-zero values. If a mask is applied to another binary or to a grayscale image of the same size, all pixels which are zero in the mask are set to zero in the output image. All others remain unchanged.

Masking can be implemented either using pixel multiplication or logical AND, the latter in general being faster.

Masking is often used to restrict a point or arithmetic operator to an area defined by the mask. We can, for example, accomplish this by first masking the desired area in the input image and processing it with the operator, then masking the original input image with the inverted mask to obtain the unprocessed area of the image and finally recombining the two partial images using image addition. An example can be seen in the worksheet on the logical AND operator. In some image processing packages, a mask can directly be defined as an optional input to a point operator, so that automatically the operator is only applied to the pixels defined by the mask .

Mathematical Morphology

The field of mathematical morphology contributes a wide range of operators to image processing, all based around a few simple mathematical concepts from set theory. The operators are particularly useful for the analysis of binary images and common usages include edge detection, noise removal, image enhancement and image segmentation.

The two most basic operations in mathematical morphology are erosion and dilation. Both of these operators take two pieces of data as input: an image to be eroded or dilated, and a structuring element (also known as a kernel). The two pieces of input data are each treated as representing sets of coordinates in a way that is slightly different for binary and grayscale

images.

For a binary image, white pixels are normally taken to represent foreground regions, while black pixels denote background. (Note that in some implementations this convention is reversed, and so it is very important to set up input images with the correct polarity for the implementation being used). Then the set of coordinates corresponding to that image is simply the set of two-dimensional Euclidean coordinates of all the foreground pixels in the image, with an origin normally taken in one of the corners so that all coordinates have positive elements.

For a grayscale image, the intensity value is taken to represent height above a base plane, so that the grayscale image represents a surface in three-dimensional Euclidean space. Figure 1 shows such a surface. Then the set of coordinates associated with this image surface is simply the set of three-dimensional Euclidean coordinates of all the points within this surface and also all points below the surface, down to the base plane. Note that even when we are only considering points with integer coordinates, this is a lot of points, so usually algorithms are employed that do not need to consider all the points.




Figure 4.12 Simple graylevel image and the corresponding surface in image space

The structuring element is already just a set of point coordinates (although it is often represented as a binary image). It differs from the input image coordinate set in that it is normally much smaller, and its coordinate origin is often not in a corner, so that some coordinate elements will have negative values. Note that in many implementations of morphological operators, the structuring element is assumed to be a particular shape (e.g. a 3×3 square) and so is hardwired into the algorithm.

Binary morphology can be seen as a special case of graylevel morphology in which the input image has only two graylevels at values 0 and 1.

Erosion and dilation work (at least conceptually) by translating the structuring element to various points in the input image, and examining the intersection between the translated kernel coordinates and the input image coordinates. For instance, in the case of erosion, the output coordinate set consists of just those points to which the origin of the structuring element can be translated, while the element still remains entirely 'within' the input image.

Virtually all other mathematical morphology operators can be defined in terms of combinations of erosion and dilation along with set operators such as intersection and union. Some of the more important are opening, closing and skeletonization.

Multi-spectral Images

A multi-spectral image is a collection of several monochrome images of the same scene, each of them taken with a different sensor. Each image is referred to as a band. A well known multi-spectral (or multi-band image) is a RGB color image, consisting of a red, a green and a blue image, each of them taken with a sensor sensitive to a different wavelength. In image processing, multi-spectral images are most commonly used for Remote Sensing applications. Satellites usually take several images from frequency bands in the visual and non-visual range. Landsat 5, for example, produces 7 band images with the wavelength of the bands being between 450 and 1250 nm.

All the standard single-band image processing operators can also be applied to multi-spectral images by processing each band separately. For example, a multi-spectral image can be edge detected by finding the edges in each band and then ORing the three edge images together. However, we would obtain more reliable edges, if we associate a pixel with an edge based on its properties in all three bands and not only in one.

To fully exploit the additional information which is contained in the multiple bands, we should consider the images as one multi-spectral image rather than as a set of monochrome graylevel images. For an image with k bands, we can then describe the brightness of each pixel as a point in a k -dimensional space represented by a vector of length k .

Special techniques exist to process multi-spectral images. For example, to classify a pixel as belonging to one particular region, its intensities in the different bands are said to form a feature vector describing its location in the k -dimensional feature space. The simplest

way to define a class is to choose a upper and lower threshold for each band, thus producing a k-dimensional 'hyper-cube' in the feature space. Only if the feature vector of a pixel points to a location within this cube, is the pixel classified as belonging to this class. A more sophisticated classification method is described in the corresponding worksheet.

The disadvantage of multi-spectral images is that, since we have to process additional data, the required computation time and memory increase significantly. However, since the speed of the hardware will increase and the costs for memory will decrease in the future, it can be expected that multi-spectral images will become more important in many fields of computer vision.

Non-linear Filtering

Suppose that an image processing operator F acting on two input images A and B produces output images C and D respectively. If the operator F is linear, then

$$F(a \times A + b \times B) = a \times C + b \times D$$

where a and b are constants. In practice, this means that each pixel in the output of a linear operator is the weighted sum of a set of pixels in the input image.

By contrast, non-linear operators are all the other operators. For example, the threshold operator is non-linear, because individually, corresponding pixels in the two images A and B may be below the threshold, whereas the pixel obtained by adding A and B may be above threshold. Similarly, an absolute value operation is non-linear:

$$|-1 + 1| \neq |-1| + |1|$$

as is the exponential operator:

$$\exp(1+1) \neq \exp(1) + \exp(1)$$

Pixels

In order for any digital computer processing to be carried out on an image, it must first be stored within the computer in a suitable form that can be manipulated by a computer program. The most practical way of doing this is to divide the image up into a collection of discrete (and usually small) cells, which are known as pixels. Most commonly, the image is

divided up into a rectangular grid of pixels, so that each pixel is itself a small rectangle. Once this has been done, each pixel is given a pixel value that represents the color of that pixel. It is assumed that the whole pixel is the same color, and so any color variation that did exist within the area of the pixel before the image was discretized is lost. However, if the area of each pixel is very small, then the discrete nature of the image is often not visible to the human eye.

Other pixel shapes and formations can be used, most notably the hexagonal grid, in which each pixel is a small hexagon. This has some advantages in image processing, including the fact that pixel connectivity is less ambiguously defined than with a square grid, but hexagonal grids are not widely used. Part of the reason is that many image capture systems (e.g. most CCD cameras and scanners) intrinsically discretize the captured image into a rectangular grid in the first instance.

Pixel Connectivity

The notation of pixel connectivity describes a relation between two or more pixels. For two pixels to be connected they have to fulfill certain conditions on the pixel brightness and spatial adjacency.

First, in order for two pixels to be considered connected, their pixel values must both be from the same set of values V . For a grayscale image, V might be any range of graylevels, e.g. $V = \{22, 23, \dots, 40\}$, for a binary image we simply have $V = \{1\}$.

To formulate the adjacency criterion for connectivity, we first introduce the notation of neighborhood. For a pixel p with the coordinates (x, y) the set of pixels given by:

$$N_4(p) = \{(x + 1, y), (x - 1, y), (x, y + 1), (x, y - 1)\}$$

is called its 4-neighbors. Its 8-neighbors are defined as

$$N_8(p) = N_4 \cup \{(x + 1, y + 1), (x + 1, y - 1), (x - 1, y + 1), (x - 1, y - 1)\}$$

From this we can infer the definition for 4- and 8-connectivity:

Two pixels p and q , both having values from a set V are 4-connected if q is from the set $N_4(p)$ and 8-connected if q is from $N_8(p)$.

General connectivity can either be based on 4- or 8-connectivity; for the following discussion we use 4-connectivity.

A pixel p is connected to a pixel q if p is 4-connected to q or if p is 4-connected to a third pixel which itself is connected to q . Or, in other words, two pixels q and p are connected

if there is a path from p and q on which each pixel is 4-connected to the next one.

This has some advantages in image processing, including the fact that pixel connectivity is less ambiguously defined than with a square grid, but hexagonal grids are not widely used. Part of the reason is that many image capture systems (e.g. most CCD cameras and scanners) intrinsically discretize the captured image into a rectangular grid in the first instance.

A set of pixels in an image which are all connected to each other is called a connected component. Finding all connected components in an image and marking each of them with a distinctive label is called connected component labeling.

An example of a binary image with two connected components which are based on 4-connectivity can be seen in Figure 1. If the connectivity were based on 8-neighbors, the two connected components would merge into one.




Figure 4.13 Two connected components based on 4-connectivity.

Pixel Values

Each of the pixels that represents an image stored inside a computer has a pixel value which describes how bright that pixel is, and/or what color it should be. In the simplest case of binary images, the pixel value is a 1-bit number indicating either foreground or background. For a grayscale images, the pixel value is a single number that represents the

brightness of the pixel. The most common pixel format is the byte image, where this number is stored as an 8-bit integer giving a range of possible values from 0 to 255. Typically zero is taken to be black, and 255 is taken to be white. Values in between make up the different shades of gray.

To represent color images, separate red, green and blue components must be specified for each pixel (assuming an RGB colorspace), and so the pixel 'value' is actually a vector of three numbers. Often the three different components are stored as three separate 'grayscale' images known as color planes (one for each of red, green and blue), which have to be recombined when displaying or processing.

Multi-spectral images can contain even more than three components for each pixel, and by extension these are stored in the same kind of way, as a vector pixel value, or as separate color planes.

The actual grayscale or color component intensities for each pixel may not actually be stored explicitly. Often, all that is stored for each pixel is an index into a colormap in which the actual intensity or colors can be looked up.

Although simple 8-bit integers or vectors of 8-bit integers are the most common sorts of pixel values used, some image formats support different types of value, for instance 32-bit signed integers or floating point values. Such values are extremely useful in image processing as they allow processing to be carried out on the image where the resulting pixel values are not necessarily 8-bit integers. If this approach is used then it is usually necessary to set up a colormap which relates particular ranges of pixel values to particular displayed colors.

Primary Colors

It is a useful fact that the huge variety of colors that can be perceived by humans can all be produced simply by adding together appropriate amounts of red, blue and green colors. These colors are known as the primary colors. Thus in most image processing applications, colors are represented by specifying separate intensity values for red, green and blue components. This representation is commonly referred to as RGB.

The primary color phenomenon results from the fact that humans have three different sorts of color receptors in their retinas which are each most sensitive to different visible light wavelengths.

The primary colors used in painting (red, yellow and blue) are different. When paints are mixed, the 'addition' of a new color paint actually subtracts wavelengths from the reflected visible light.

RGB and Colorspaces

A color perceived by the human eye can be defined by a linear combination of the three primary colors red, green and blue. These three colors form the basis for the RGB-colorspace. Hence, each perceivable color can be defined by a vector in the three-dimensional colorspace. The intensity is given by the length of the vector, and the actual color by the two angles describing the orientation of the vector in the colorspace.

The RGB-space can also be transformed into other coordinate systems, which might be more useful for some applications. One common basis for the color space is IHS. In this coordinate system, a color is described by its intensity, hue (average wavelength) and saturation (the amount of white in the color). This color space makes it easier to directly derive the intensity and color of perceived light and is therefore more likely to be used by human beings.

Spatial Domain

For simplicity, assume that the image I being considered is formed by projection from scene S (which might be a two- or three-dimensional scene, etc.).

The spatial domain is the normal image space, in which a change in position in I directly projects to a change in position in S . Distances in I (in pixels) correspond to real distances (e.g. in meters) in S .

This concept is used most often when discussing the frequency with which image values change, that is, over how many pixels does a cycle of periodically repeating intensity variations occur. One would refer to the number of pixels over which a pattern repeats (its periodicity) in the spatial domain.

In most cases, the Fourier Transform will be used to convert images from the spatial domain into the frequency domain.

A related term used in this context is spatial frequency, which refers to the (inverse of the) periodicity with which the image intensity values change. Image features with high spatial frequency (such as edges) are those that change greatly in intensity over short image

distances.

Another term used in this context is spatial derivative, which refers to how much the image intensity values change per change in image position.

Structuring Elements

The field of mathematical morphology provides a number of important image processing operations, including erosion, dilation, opening and closing. All these morphological operators take two pieces of data as input. One is the input image, which may be either binary or grayscale for most of the operators. The other is the structuring element. It is this that determines the precise details of the effect of the operator on the image.

The structuring element is sometimes called the kernel, but we reserve that term for the similar objects used in convolutions.

The structuring element consists of a pattern specified as the coordinates of a number of discrete points relative to some origin. Normally cartesian coordinates are used and so a convenient way of representing the element is as a small image on a rectangular grid. Figure 1 shows a number of different structuring elements of various sizes. In each case the origin is marked by a ring around that point. The origin does not have to be in the center of the structuring element, but often it is. As suggested by the figure, structuring elements that fit into a 3×3 grid with its origin at the center are the most commonly seen type.




Figure 4.14 Some example structuring elements.

Note that each point in the structuring element may have a value. In the simplest structuring elements used with binary images for operations such as erosion, the elements only have one value, conveniently represented as a one. More complicated elements, such as

those used with thinning or grayscale morphological operations, may have other pixel values.

An important point to note is that although a rectangular grid is used to represent the structuring element, not every point in that grid is part of the structuring element in general. Hence the elements shown in Figure 1 contain some blanks. In many texts, these blanks are represented as zeros, but this can be confusing and so we avoid it here.

When a morphological operation is carried out, the origin of the structuring element is typically translated to each pixel position in the image in turn, and then the points within the translated structuring element are compared with the underlying image pixel values. The details of this comparison, and the effect of the outcome depend on which morphological operator is being used.

Wrapping and Saturation

If an image is represented in a byte or integer pixel format, the maximum pixel value is limited by the number of bits used for the representation, e.g. the pixel values of a 8-bit image are limited to 255.

However, many image processing operations produce output values which are likely to exceed the given maximum value. In such cases, we have to decide how to handle this pixel overflow.

One possibility is to wrap around the overflowing pixel values. This means that if a value is greater than the possible maximum, we subtract the pixel value range so that the value starts again from the possible minimum value. Figure 1 shows the mapping function for wrapping the output values of some operation into an 8-bit format.




Figure 4.15 Mapping function for wrapping the pixel values of an 8-bit image.

Another possibility is to set all overflowing pixels to the maximum possible values --- an effect known as saturation. The corresponding mapping function for an 8-bit image can be seen in Figure 4.6

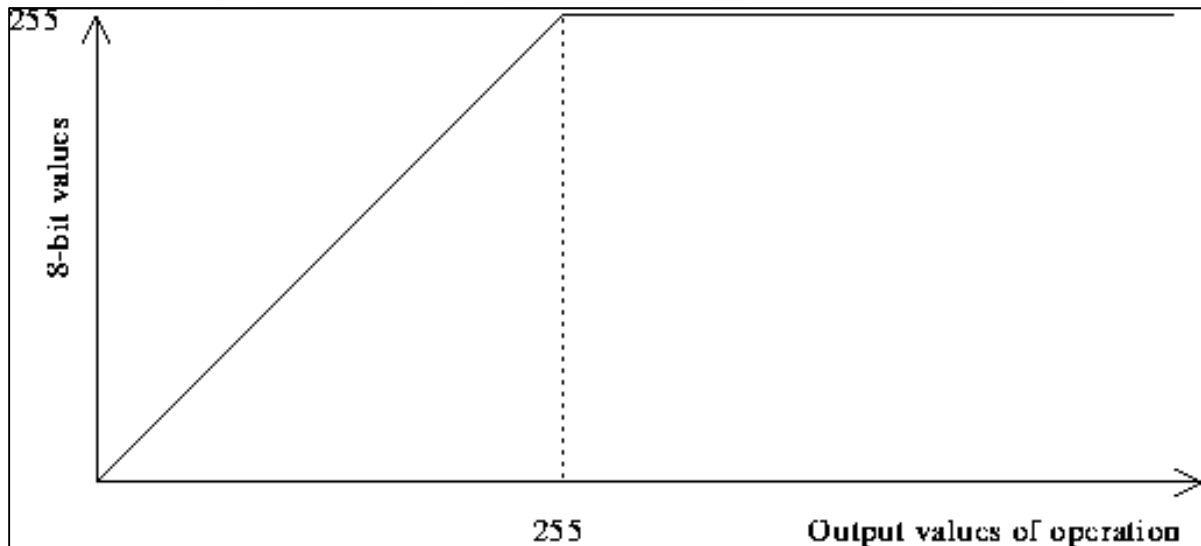


Figure 4.16 Mapping function for saturating an 8-bit image.

If only a few pixels in the image exceed the maximum value it is often better to apply the latter technique, especially if we use the image for display purposes. However, by setting all overflowing pixels to the same value we lose an essential amount of information. In the worst case, when all pixels exceed the maximum value, this would lead to an image of constant pixel values. Wrapping around overflowing pixel retains the differences between values. On the other hand, it might cause the problem that pixel values passing the maximum 'jump' from the maximum to the minimum value. Examples for both techniques can be seen in the worksheets of various point operators.

If possible, it is easiest to change the image format, for example to float format, so that all pixel values can be represented. However, we should keep in mind that this implies an increase in processing time and memory.

Input Design

The input design is the link between the information system and the user. It comprises the developing specification and procedures for data preparation and those steps are necessary to put transaction data in to a usable form for processing can be achieved by inspecting the computer to read data from a written or printed document or it can occur by having people keying the data directly into the system. The design of input focuses on controlling the amount of input required, controlling the errors, avoiding delay, avoiding extra steps and keeping the process simple. The input is designed in such a way so that it

provides security and ease of use with retaining the privacy. Input Design considered the following things:

- What data should be given as input?
- How the data should be arranged or coded?
- The dialog to guide the operating personnel in providing input.
- Methods for preparing input validations and steps to follow when error occur.

OUTPUT DESIGN

A quality output is one, which meets the requirements of the end user and presents the information clearly. In any system results of processing are communicated to the users and to other system through outputs. In output design it is determined how the information is to be displaced for immediate need and also the hard copy output. It is the most important and direct source information to the user. Efficient and intelligent output design improves the system's relationship to help user decision-making. The output form of an information system should accomplish one or more of the following objectives.

- Convey information about past activities, current status or projections of the Future.
- Signal important events, opportunities, problems, or warnings.
- Trigger an action.
- Confirm an action.

CHAPTER - 5

TECHNOLOGY AND TOOLS

SOFTWARE MODULES

- **IMAGE PROCESSING TOOL: OPENCV**
- **PROGRAM: PYTHON**

5.1. PYTHON (PROGRAMMING LANGUAGE)

Python is a widely used high-level programming language for general-purpose programming, created by Guido van Rossum and first released in 1991. An interpreted language, Python has a design philosophy which emphasizes code readability (notably using whitespace indentation to delimit code blocks rather than curly brackets or keywords), and a syntax which allows programmers to express concepts in fewer lines of code than possible in languages such as C++ or Java. The language provides constructs intended to enable writing clear programs on both a small and large scale. Python features a dynamic type system and automatic memory management and supports multiple programming paradigms, including object-oriented, imperative, functional programming, and procedural styles. It has a large and comprehensive standard library. Python interpreters are available for many operating systems, allowing Python code to run on a wide variety of systems. CPython, the reference implementation of Python, is open source software and has a community-based development model, as do nearly all of its variant implementations. CPython is managed by the non-profit Python Software Foundation

5.2 HISTORY

Python was conceived in the late 1980s, and its implementation began in December 1989 [28] by Guido van Rossum at Centrum Wiskunde & Informatica (CWI) in the Netherlands as a successor to the ABC language (itself inspired by SETL) capable of exception handling and interfacing with the operating system Amoeba. Van Rossum is Python's principal author, and his continuing central role in deciding the direction of Python is reflected in the title given to him by the Python community, benevolent dictator for life (BDFL).

About the origin of Python, Van Rossum wrote in 1996: Over six years ago, in December 1989, I was looking for a "hobby" programming project that would keep me occupied during the week around Christmas. My office ... would be closed, but I had a home computer, and not much else on my hands. I decided to write an interpreter for the new

scripting language I had been thinking about lately: a descendant of ABC that would appeal to Unix/C hackers. I chose Python as a working title for the project, being in a slightly irreverent mood (and a big fan of Monty Python's Flying Circus). Python 2.0 was released on 16 October 2000 and had many major new features, including a cycle-detecting garbage collector and support for Unicode. With this release the development process was changed and became more transparent and community-backed. Python 3.0 (which early in its development was commonly referred to as Python 3000 or py3k), a major, backwards-incompatible release, was released on 3 December 2008 after a long period of testing.

Many of its major features have been backported to the backwards-compatible Python 2.6.x and 2.7.x version series. The End Of Life date (EOL, sunset date) for Python 2.7 was initially set at 2015, then postponed to 2020 out of concern that a large body of existing code cannot easily be forward-ported to Python 3. In January 2017, Google announced work on a Python 2.7 to go transcompiler, which The Register speculated was in response to Python 2.7's planned end-of-life but Google cited performance under concurrent workloads as their only motivation.

Features and philosophy Python is a multi-paradigm programming language: object-oriented programming and structured programming are fully supported, and many language features support functional programming and aspect-oriented programming (including by metaprogramming and metaobjects (magic methods)). Many other 6/13/2017 Python (programming language) paradigms are supported via extensions, including design by contract and logic programming. Python uses dynamic typing and a mix of reference counting and a cycle-detecting garbage collector for memory management. An important feature of Python is dynamic name resolution (late binding), which binds method and variable names during program execution.

The design of Python offers some support for functional programming in the Lisp tradition. The language has `map()`, `reduce()` and `filter()` functions; list comprehensions, dictionaries, and sets; and generator expressions. The standard library has two modules (`itertools` and `functools`) that implement functional tools borrowed from Haskell and Standard ML.

The core philosophy of the language is summarized by the document The Zen of Python (PEP 20), which includes aphorisms such as: Beautiful is better than ugly Explicit is better than implicit Simple is better than complex Complex is better than complicated Readability counts Rather than requiring all desired functionality to be built into the language's core, Python was designed to be highly extensible. Python can also be embedded

in existing applications that need a programmable interface.

This design of a small core language with a large standard library and an easily extensible interpreter was intended by Van Rossum from the start because of his frustrations with ABC, which espoused the opposite mindset. While offering choice in coding methodology, the Python philosophy rejects exuberant syntax, such as in Perl, in favor of a sparser, less-cluttered grammar.

As Alex Martelli put it: "To describe something as clever is not considered a compliment in the Python culture." Python's philosophy rejects the Perl "there is more than one way to do it" approach to language design in favor of "there should be one—and preferably only one—obvious way to do it". Python's developers strive to avoid premature optimization, and moreover, reject patches to non-critical parts of C Python that would offer a marginal increase in speed at the cost of clarity. When speed is important, a Python programmer can move time-critical functions to extension modules written in languages such as C, or try using PyPy, a just-in-time compiler. Cython is also available, which translates a Python script into C and makes direct C-level API calls into the Python interpreter. An important goal of Python's developers is making it fun to use. This is reflected in the origin of the name, which comes from Monty Python, and in an occasionally playful approach to tutorials and reference materials, such as using examples that refer to spam and eggs instead of the standard foo and bar. A common neologism in the Python community is *pythonic*, which can have a wide range of meanings related to program style.

To say that code is *pythonic* is to say that it uses Python idioms well, that it is natural or shows fluency in the language, that it conforms with Python's minimalist philosophy and emphasis on readability. In contrast, code that is difficult to understand or reads like a rough transcription from another programming language is called *unpythonic*. Users and admirers of Python, especially those considered knowledgeable or experienced, are often referred to as *Pythonists*, *Pythonistas*, and *Pythoneers*. The third slice parameter, called *step* or *stride*, allows elements to be skipped and reversed. Slice indexes may be omitted, for example `a[:]` returns a copy of the entire list. Each element of a slice is a shallow copy. In Python, a distinction between expressions and statements is rigidly enforced, in contrast to languages such as Common Lisp, Scheme, or Ruby. This leads to duplicating some functionality.

For example: List comprehensions vs. for-loops Conditional expressions vs. if blocks
The `eval()` vs. `exec()` built-in functions (in Python 2, `exec` is a statement); the former is for expressions, the latter is for statements. 6/13/2017 Python (programming language) - Statements cannot be a part of an expression, so list and other comprehensions or lambda

expressions, all being expressions, cannot contain statements.

A particular case of this is that an assignment statement such as `a = 1` cannot form part of the conditional expression of a conditional statement. This has the advantage of avoiding a classic C error of mistaking an assignment operator `=` for an equality operator `==` in conditions: `if (c = 1) { ... }` is syntactically valid (but probably unintended) C code but `if c = 1: ...` causes a syntax error in Python. Methods Methods on objects are functions attached to the object's class; the syntax `instance.method(argument)` is, for normal methods and functions, syntactic sugar for `Class.method(instance, argument)`. Python methods have an explicit self parameter to access instance data, in contrast to the implicit self (or this) in some other object-oriented programming languages (e.g., C++, Java, Objective-C, or Ruby). Typing Python uses duck typing and has typed objects but untyped variable names. Type constraints are not checked at compile time; rather, operations on an object may fail, signifying that the given object is not of a suitable type. Despite being dynamically typed, Python is strongly typed, forbidding operations that are not well-defined (for example, adding a number to a string) rather than silently attempting to make sense of them. Python allows programmers to define their own types using classes, which are most often used for objectoriented programming.

New instances of classes are constructed by calling the class (for example, `SpamClass()` or `EggsClass()`), and the classes are instances of the metaclass type (itself an instance of itself), allowing meta programming and reflection. Before version 3.0, Python had two kinds of classes: old-style and new-style. The syntax of both styles is the same, the difference being whether the class object is inherited from, directly or indirectly (all new-style classes inherit from `object` and are instances of `type`). In versions of Python 2 from Python 2.2 onwards, both kinds of classes can be used. Old-style classes were eliminated in Python 3.0. The long term plan is to support gradual typing and as of Python 3.5, the syntax of the language allows specifying static types but they are not checked in the default implementation, CPython. An experimental optional static type checker named `mypy` supports compile-time type checking.

Mathematics Python has the usual C arithmetic operators (`+`, `-`, `*`, `/`, `%`). It also has `**` for exponentiation, e.g. `5**3 == 125` and `9**0.5 == 3.0`, and a new matrix multiply `@` operator is included in version 3.5. [71] Additionally, it has a unary operator (`~`), which essentially inverts all the bytes of its one argument. For integers, this means `~x == -x-1`. Other operators include bitwise shift operators `x << y`, which shifts `x` to the left `y` places, the same as `x*(2**y)`, and `x >> y`, which shifts `x` to the right `y` places, the same as `x/(2**y)`. [73] The

behavior of division has changed significantly over time: Python 2.1 and earlier use the C division behavior. The `/` operator is integer division if both operands are integers, and floating-point division otherwise. Integer division rounds towards 0, e.g. $7 / 3 == 2$ and $-7 / 3 == -2$. Python 2.2 changes integer division to round towards negative infinity, e.g. $7 / 3 == 2$ and $-7 / 3 == -3$. The floor division `//` operator is introduced. So $7 // 3 == 2$, $-7 // 3 == -3$, $7.5 // 3 == 2.0$ and $-7.5 // 3 == -3.0$.
[https://en.wikipedia.org/wiki/Python_\(programming_language\)](https://en.wikipedia.org/wiki/Python_(programming_language)) 8/19

Adding `from_futur import division` causes a module to use Python 3.0 rules for division (see next). Python 3.0 changes `/` to be always floating-point division. In Python terms, the pre-3.0 `/` is classic division, the version-3.0 `/` is real division, and `//` is floor division. Rounding towards negative infinity, though different from most languages, adds consistency. For instance, it means that the equation $(a+b) // b == a // b + 1$ is always true. It also means that the equation $b * (a // b) + a \% b == a$ is valid for both positive and negative values of a . However, maintaining the validity of this equation means that while the result of $a \% b$ is, as expected, in the half-open interval $[0, b)$, where b is a positive integer, it has to lie in the interval $(b, 0]$ when b is negative. Python provides a `round` function for rounding a float to the nearest integer. For tie-breaking, versions before 3 use round-away-from-zero: `round(0.5)` is 1.0, `round(-0.5)` is -1.0. Python 3 uses round to even: `round(1.5)` is 2, `round(2.5)` is 2. Python allows boolean expressions with multiple equality relations in a manner that is consistent with general use in mathematics.

For example, the expression $a < b < c$ tests whether a is less than b and b is less than c . C-derived languages interpret this expression differently: in C, the expression would first evaluate $a < b$, resulting in 0 or 1, and that result would then be compared with c . Python has extensive built-in support for arbitrary precision arithmetic. Integers are transparently switched from the machine-supported maximum fixed-precision (usually 32 or 64 bits), belonging to the python type `int`, to arbitrary precision, belonging to the python type `long`, where needed. The latter have an "L" suffix in their textual representation. (In Python 3, the distinction between the `int` and `long` types was eliminated; this behavior is now entirely contained by the `int` class.) The `Decimal` type/class in module `decimal` (since version 2.4) provides decimal floating point numbers to arbitrary precision and several rounding modes. The `Fraction` type in module `fractions` (since version 2.6) provides arbitrary precision for rational numbers.

Due to Python's extensive mathematics library, and the third-party library NumPy which further extends the native capabilities, it is frequently used as a scientific scripting

language to aid in problems such as numerical data processing and manipulation. Libraries Python has a large standard library, commonly cited as one of Python's greatest strengths, providing tools suited to many tasks. This is deliberate and has been described as a "batteries included" Python philosophy.

5.3 OPENCV

OpenCV (*Open source computer vision*) is a library of programming functions mainly aimed at real-time computer vision. Originally developed by Intel, it was later supported by Willow Garage then Itseez (which was later acquired by Intel). The library is cross-platform and free for use under the open-source BSD license.

OpenCV supports the deep learning frameworks TensorFlow, Torch/PyTorch and Caffe

5.3.1 OpenCV HISTORY

Officially launched in 1999 the OpenCV project was initially an Intel Research initiative to advance CPU-intensive applications, part of a series of projects including real-time ray tracing and 3D display walls. The main contributors to the project included a number of optimization experts in Intel Russia, as well as Intel's Performance Library Team. In the early days of OpenCV, the goals of the project were described as:

- Advance vision research by providing not only open but also optimized code for basic vision infrastructure. No more reinventing the wheel.
- Disseminate vision knowledge by providing a common infrastructure that developers could build on, so that code would be more readily readable and transferable.
- Advance vision-based commercial applications by making portable, performance-optimized code available for free – with a license that did not require code to be open or free itself.

The first alpha version of OpenCV was released to the public at the IEEE Conference on Computer Vision and Pattern Recognition in 2000, and five betas were released between 2001 and 2005. The first 1.0 version was released in 2006. A version 1.1 "pre-release" was released in October 2008.

The second major release of the OpenCV was in October 2009. OpenCV 2 includes major changes to the C++ interface, aiming at easier, more type-safe patterns, new functions, and better implementations for existing ones in terms of performance (especially on multi-core systems). Official releases now occur every six months¹ and development is now done by an independent Russian team supported by commercial corporations.

In August 2012, support for OpenCV was taken over by a non-profit foundation OpenCV.org, which maintains a developer and user site.

On May 2016, Intel signed an agreement to acquire Itseez, a leading developer of OpenCV.

5.3.2 APPLICATIONS

OpenCV's application areas include:

- 2D and 3D feature toolkits
- Egomotion estimation
- Facial recognition system
- Gesture recognition
- Human–computer interaction (HCI)
- Mobile robotics
- Motion understanding
- Object identification
- Segmentation and recognition
- Stereopsis stereo vision: depth perception from 2 cameras
- Structure from motion (SFM)
- Motion tracking
- Augmented reality

To support some of the above areas, OpenCV includes a statistical machine learning library that contains:

- Boosting
- Decision tree learning
- Gradient boosting trees

- Expectation-maximization algorithm
- k-nearest neighbor algorithm
- Naive Bayes classifier
- Artificial neural networks
- Random forest
- Support vector machine (SVM)
- Deep neural networks (DNN)

5.4 PROGRAMMING LANGUAGE

OpenCV is written in C++ and its primary interface is in C++, but it still retains a less comprehensive though extensive older C interface. There are bindings in Python, Java and MATLAB/OCTAVE. The API for these interfaces can be found in the online documentation. Wrappers in other languages such as C#, Perl, Ch, Haskell, and Ruby have been developed to encourage adoption by a wider audience.

Since version 3.4, OpenCV.js is a JavaScript binding for selected subset of OpenCV functions for the web platform.

All of the new developments and algorithms in OpenCV are now developed in the C++ interface.

HARDWARE ACCELERATION

If the library finds Intel's Integrated Performance Primitives on the system, it will use these proprietary optimized routines to accelerate itself. A CUDA-based GPU interface has been in progress since September 2010. An OpenCL-based GPU interface has been in progress since October 2012, documentation for version 2.4.13.3 can be found at docs.opencv.org.

5.4.1 OS SUPPORT

OpenCV runs on the following desktop operating systems: Windows, Linux, macOS, FreeBSD, NetBSD, OpenBSD. OpenCV runs on the following mobile operating systems: Android, iOS, Maemo, BlackBerry 10. The user can get official releases from SourceForge or take the latest sources from GitHub. OpenCV uses CMake.

Conclusion:

In order to detect a driver's drowsiness, facial features, eyes and mouth were identified on the video of an individual driving. Convolutional neural network was implemented to classify eyes as open or closed. Drowsiness was determined on the basis of frequency of closed eyes. Using OpenCV and Dlib in Python, frequency of yawning was examined. An alarm was set to ring after the detection to alert the driver. There will be limitations concerning the detection of drivers' conditions and facial expressions due to factors like darkness, light reflection, obstructions by drivers' hands and wearing of sunglasses. Convolutional neural gives better performance and facial extraction method accompanies it, as an additional drowsiness detection technique which is often used with other facial extraction methods.

CHAPTER – 6

CONCLUSION AND FUTURE ENHANCEMENT

Driver drowsiness detection using Convolutional Neural Networks (CNNs) is a critical application in ensuring road safety. CNNs excel in image recognition tasks, making them ideal for analyzing facial features indicative of drowsiness. The process typically involves capturing real-time images or video frames of the driver's face using a camera installed in the vehicle. These images are then fed into the CNN model, which learns to detect patterns associated with drowsiness, such as eye closure, head nodding, or changes in facial expressions.

The CNN architecture consists of multiple layers, including convolutional layers for feature extraction and pooling layers for dimensionality reduction. These layers enable the model to automatically learn relevant features from the input images without manual feature engineering. Additionally, techniques like transfer learning can be employed to leverage pre-trained CNN models, further enhancing performance with limited training data.

Once trained, the CNN model can continuously analyze incoming images or frames in real-time, providing timely alerts to the driver or triggering automated safety mechanisms if drowsiness is detected. Such systems play a crucial role in preventing accidents caused by driver fatigue, potentially saving lives on the road. Overall, CNN-based drowsiness detection systems offer a robust and efficient solution for enhancing driver safety through advanced image analysis techniques.

6.1 Conclusion

In conclusion, driver drowsiness detection systems represent a critical advancement in automotive safety technology, aiming to mitigate the risks associated with fatigued driving and prevent potentially devastating accidents. Through the integration of various sensors, algorithms, and real-time monitoring techniques, these systems offer a proactive approach to addressing driver fatigue, a significant contributor to road accidents worldwide.

The implementation of computer vision techniques, particularly utilizing Convolutional Neural Networks (CNNs), has proven instrumental in accurately detecting signs of drowsiness by analyzing facial expressions, eye movements, and head poses. These CNN-based systems can effectively distinguish between alert and drowsy states, providing timely alerts or interventions to prevent accidents. Moreover, the utilization of additional sensors such as steering wheel sensors, heart rate monitors, and lane departure detection

systems further enhances the robustness and reliability of drowsiness detection systems, enabling them to adapt to diverse driving conditions and individual differences in behavior.

Real-time processing capabilities allow these systems to continuously monitor driver behavior and promptly respond to signs of drowsiness, ensuring optimal safety on the road. By integrating advanced algorithms with edge computing technology, drowsiness detection systems can operate with low latency, enabling rapid decision-making without relying on cloud-based processing, thus ensuring uninterrupted performance even in areas with limited connectivity.

Furthermore, the evolution of driver drowsiness detection systems extends beyond mere detection to encompass proactive interventions aimed at preventing fatigue-related accidents altogether. This includes adaptive cruise control systems, lane-keeping assistance, and autonomous driving features that can intervene when a driver is deemed unfit to operate the vehicle safely. These interventions not only enhance safety but also promote user comfort and convenience, contributing to a more enjoyable driving experience.

The widespread adoption of driver drowsiness detection systems has the potential to significantly reduce the incidence of fatigue-related accidents, saving lives, and reducing economic costs associated with road traffic accidents. Governments, regulatory bodies, and automotive manufacturers play a crucial role in promoting the adoption of these technologies through the implementation of regulations, incentives, and public awareness campaigns highlighting the importance of driver safety.

However, challenges remain, including the need for further research and development to improve the accuracy and reliability of drowsiness detection algorithms, particularly in diverse environmental conditions and for drivers with varying physiological characteristics. Additionally, addressing privacy concerns related to the collection and processing of driver data is essential to ensure widespread acceptance and adoption of these systems.

In conclusion, driver drowsiness detection systems represent a promising advancement in automotive safety technology, offering the potential to save lives and prevent accidents by proactively identifying and addressing driver fatigue. Through the integration of advanced sensors, algorithms, and real-time monitoring capabilities, these systems have the capacity to revolutionize road safety and usher in a new era of safer, more efficient transportation.

6.2 Future Enhancement

Future enhancements for driver drowsiness detection systems will likely focus on improving accuracy, real-time performance, adaptability to diverse driving conditions, and user experience. One avenue for enhancement involves the integration of multimodal sensor data, including not only facial images but also physiological signals such as heart rate variability (HRV), electroencephalogram (EEG) data, and steering wheel dynamics. Combining these modalities can provide a more comprehensive understanding of the driver's cognitive and physiological state, leading to more robust drowsiness detection.

Another area of advancement is the development of deep learning architectures specifically tailored for drowsiness detection. While Convolutional Neural Networks (CNNs) have shown promise in analyzing facial features, there is potential for further innovation in the design of neural network architectures to capture subtle cues of drowsiness from facial expressions, eye movements, and head poses. Recurrent Neural Networks (RNNs) and attention mechanisms can be explored to model temporal dependencies in driver behavior, enabling more accurate prediction of drowsiness onset.

Furthermore, the deployment of edge computing and onboard processing capabilities can facilitate real-time analysis of sensor data directly within the vehicle, reducing latency and enabling faster response times to drowsiness events. This involves optimizing algorithms for efficiency and resource-constrained environments while ensuring data privacy and security.

In addition to technical improvements, future enhancements should also consider the user experience aspect of drowsiness detection systems. This includes designing intuitive user interfaces, providing timely and context-aware alerts to the driver, and integrating adaptive feedback mechanisms to mitigate false alarms and adapt to individual driving behaviors. Human-centered design principles should guide the development process to ensure that drowsiness detection systems are user-friendly and non-intrusive, ultimately enhancing driver safety and acceptance.

Moreover, advancements in sensor technologies, such as the integration of non-intrusive physiological sensors into steering wheels, seats, or wearable devices, hold promise for unobtrusive monitoring of driver drowsiness without requiring explicit user interaction. These sensors can capture physiological signals indicative of drowsiness, such as changes in heart rate, respiration rate, and skin conductance, providing additional insights for drowsiness

detection algorithms.

Additionally, leveraging emerging technologies such as affective computing and affective sensing can enable drowsiness detection systems to infer the driver's emotional state and cognitive load from facial expressions, speech patterns, and behavioral cues. By considering the driver's emotional and cognitive context, these systems can enhance the accuracy and robustness of drowsiness detection algorithms, leading to more personalized and adaptive interventions.

Overall, future enhancements for driver drowsiness detection systems should adopt a holistic approach that integrates advances in sensor technologies, machine learning algorithms, edge computing, and user experience design to create intelligent, adaptive, and user-friendly systems that prioritize driver safety and well-being on the road.

APPENDIX

Driver drowsiness detection using Convolutional Neural Networks (CNNs) is a critical application in ensuring road safety. CNNs excel in image recognition tasks, making them ideal for analyzing facial features indicative of drowsiness. The process typically involves capturing real-time images or video frames of the driver's face using a camera installed in the vehicle. These images are then fed into the CNN model, which learns to detect patterns associated with drowsiness, such as eye closure, head nodding, or changes in facial expressions.

The CNN architecture consists of multiple layers, including convolutional layers for feature extraction and pooling layers for dimensionality reduction. These layers enable the model to automatically learn relevant features from the input images without manual feature engineering. Additionally, techniques like transfer learning can be employed to leverage pre-trained CNN models, further enhancing performance with limited training data.

Once trained, the CNN model can continuously analyze incoming images or frames in real-time, providing timely alerts to the driver or triggering automated safety mechanisms if drowsiness is detected. Such systems play a crucial role in preventing accidents caused by driver fatigue, potentially saving lives on the road. Overall, CNN-based drowsiness detection systems offer a robust and efficient solution for enhancing driver safety through advanced image analysis techniques.

7.1 REQUIREMENTS

- numpy==1.15.2
- dlib==19.16.0
- pygame==1.9.4
- imutils==0.5.1
- opencv_python==3.4.3.18
- scipy==1.1.0

7.2 SOURCE CODE

```
#Import necessary libraries
```

```
from scipy.spatial import distance
```

```

from imutils import face_utils

import numpy as np

import pygame #For playing sound

import time

import dlib

import cv2

#Initialize Pygame and load music

pygame.mixer.init()

pygame.mixer.music.load('audio/alert.wav')

#Minimum threshold of eye aspect ratio below which alarm is triggered

EYE_ASPECT_RATIO_THRESHOLD = 0.3

#Minimum consecutive frames for which eye ratio is below threshold for alarm
to be triggered

EYE_ASPECT_RATIO_CONSEC_FRAMES = 50
#Counts no. of consecutive frames below threshold value

COUNTER = 0

#Load face cascade which will be used to draw a rectangle around detected
faces.

face_cascade =
cv2.CascadeClassifier("haarcascades/haarcascade_frontalface_default.xml")

```

```

#This function calculates and return eye aspect ratio

def eye_aspect_ratio(eye):

    A = distance.euclidean(eye[1], eye[5])

    B = distance.euclidean(eye[2], eye[4])

    C = distance.euclidean(eye[0], eye[3])

    ear = (A+B) / (2*C)

    return ear

#Load face detector and predictor, uses dlib shape predictor file

detector = dlib.get_frontal_face_detector()

predictor = dlib.shape_predictor('shape_predictor_68_face_landmarks.dat')

#Extract indexes of facial landmarks for the left and right eye
(lStart, lEnd) = face_utils.FACIAL_LANDMARKS_IDXS['left_eye']

(rStart, rEnd) = face_utils.FACIAL_LANDMARKS_IDXS['right_eye']

#Start webcam video capture

video_capture = cv2.VideoCapture(0)

#Give some time for camera to initialize(not required)

time.sleep(2)

while(True):

    #Read each frame and flip it, and convert to grayscale

    ret, frame = video_capture.read()

    frame = cv2.flip(frame,1)

    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)

```

```

#Detect facial points through detector function

faces = detector(gray, 0)

#Detect faces through haarcascade_frontalface_default.xml

face_rectangle = face_cascade.detectMultiScale(gray, 1.3, 5)

#Draw rectangle around each face detected
for (x,y,w,h) in face_rectangle:

    cv2.rectangle(frame,(x,y),(x+w,y+h),(255,0,0),2)

#Detect facial points

for face in faces:

    shape = predictor(gray, face)

    shape = face_utils.shape_to_np(shape)

    #Get array of coordinates of leftEye and rightEye

    leftEye = shape[lStart:lEnd]

    rightEye = shape[rStart:rEnd]

    #Calculate aspect ratio of both eyes

    leftEyeAspectRatio = eye_aspect_ratio(leftEye)

    rightEyeAspectRatio = eye_aspect_ratio(rightEye)

    eyeAspectRatio = (leftEyeAspectRatio + rightEyeAspectRatio) / 2

    #Use hull to remove convex contour discrepancies and draw eye shape
around eyes

    leftEyeHull = cv2.convexHull(leftEye)
    rightEyeHull = cv2.convexHull(rightEye)
    cv2.drawContours(frame,
    [leftEyeHull], -1, (0, 255, 0), 1)
    cv2.drawContours(frame, [rightEyeHull], -1, (0, 255, 0), 1)

```



```

#Detect if eye aspect ratio is less than threshold

if(eyeAspectRatio < EYE_ASPECT_RATIO_THRESHOLD):

    COUNTER += 1

    #If no. of frames is greater than threshold frames,

    if COUNTER >= EYE_ASPECT_RATIO_CONSEC_FRAMES:

        pygame.mixer.music.play(-1)

        cv2.putText(frame, "You are Drowsy", (150,200),
cv2.FONT_HERSHEY_SIMPLEX, 1.5, (0,0,255), 2)

    else:

        pygame.mixer.music.stop()

        COUNTER = 0


#Show video feed

cv2.imshow('Video', frame)

if(cv2.waitKey(1) & 0xFF == ord('q')):

    break

#Finally when video capture is over, release the video capture and
destroyAllWindows

video_capture.release()

cv2.destroyAllWindows()

```

7.3 Sample Output




Figure 7.1 : Awake Result




Figure 7.2 : Drowsiness Result

REFERENCES

1. VU Maheswari, R Aluvalu, MVVP Kantipudi, KK Chennam, K Kotecha, JR Saini. (2022). Driver Drowsiness Prediction Based on Multiple Aspects Using Image Processing Techniques. In 2022 IEEE International Conference on Imaging Systems and Techniques (IST)
2. KF Lee, XZ Chen, CW Yu, KY Chin, YC Wang, CY Hsiao, YL Chen. (2022). An Intelligent Driving Assistance System Based on Lightweight Deep Learning models. In 2022 IEEE International Conference on Lightweight Deep Learning models
3. A Altameem, A Kumar, RC Poonia, S Kumar, AKJ Saudagar. (2021). Early Identification and Detection of Driver Drowsiness by Hybrid Machine Learning. In 2021 IEEE International Conference on Hybrid Machine Learning.
4. W Deng, R Wu. (2019). Real-Time Driver-Drowsiness Detection System Using Facial Features. In 2019 IEEE International Conference on Facial Features..
5. Yeresime Suresh; Rashi Khandelwal; Matam Nikitha; Mohammed Fayaz; Vinaya Soudhri. (2021). Driver Drowsiness Detection Using Deep Learning. In 2021 IEEE International Conference on Smart Electronics and Communication (ICOSEC)
6. N Prasath, J Sreemathy, P Vigneshwaran. (2022). Driver Drowsiness Detection Using Machine Learning Algorithm. In 2021 IEEE International Conference on Advanced Computing and Communication Systems (ICACCS)
7. Authors: IR Adochiei, OI Știrbu, NI Adochiei, M Pericle-Gabriel, CM Larco, SM Mustata, D Costin. (2020). Drivers' Drowsiness Detection and Warning Systems for Critical Infrastructures. 2020 International Conference on e-Health and Bioengineering (EHB)
8. S Ananthi, R Sathya, K Vaidehi, G Vijaya. (2023). Drivers Drowsiness Detection using Image Processing and I-Ear Techniques. 2023 7th International Conference on Intelligent Computing and Control Systems (ICICCS)

9. Miankuan Zhu; Haobo Li; Jiangfan Chen; Mitsuhiro Kamezaki; Zutao Zhang; ZexiHua; Shigeki Sugano.(2021). EEG-based System Using Deep Learning and Attention Mechanism for Driver Drowsiness Detection. In2021 IEEE Intelligent Vehicles Symposium Workshops (IV Workshops)
10. Qi Shen, Shengjie Zhao, Rongqing Zhang, Bin Zhang. (2020). Robust Two-Stream Multi-Features Network for Driver Drowsiness Detection. In 2020 IEEE International Conference on Multi Features Network
11. Chao Zhou, Pengfei Zhao, Bahareh Nakisa, Mohammad Naim Rastgoo, Ramesh Jain, Wen Gao. (2021). Driver stress detection via multimodal fusion using attention-based CNN-LSTM. In 2021 IEEE International Conference

PUBLICATIONS

Dr.P.Vasuki, B Deepak, B Naga Sudhakar, Ch Rakesh, D Tirumalesh “ CNN BASED STAY AWAKE STAY ALIVE SYSTEM FOR ACCIDENT PREVENTION ” International conference of “Artificial Intelligence in Nanomaterials Engineering, Environmental And Health care Applications (ICANEHA-2024) through online held at Asia Pacific University, Kuala Lumpur, Malaysia during 27th – 28th February 2024”